



# Kubernetes By Example

@rorypreddy



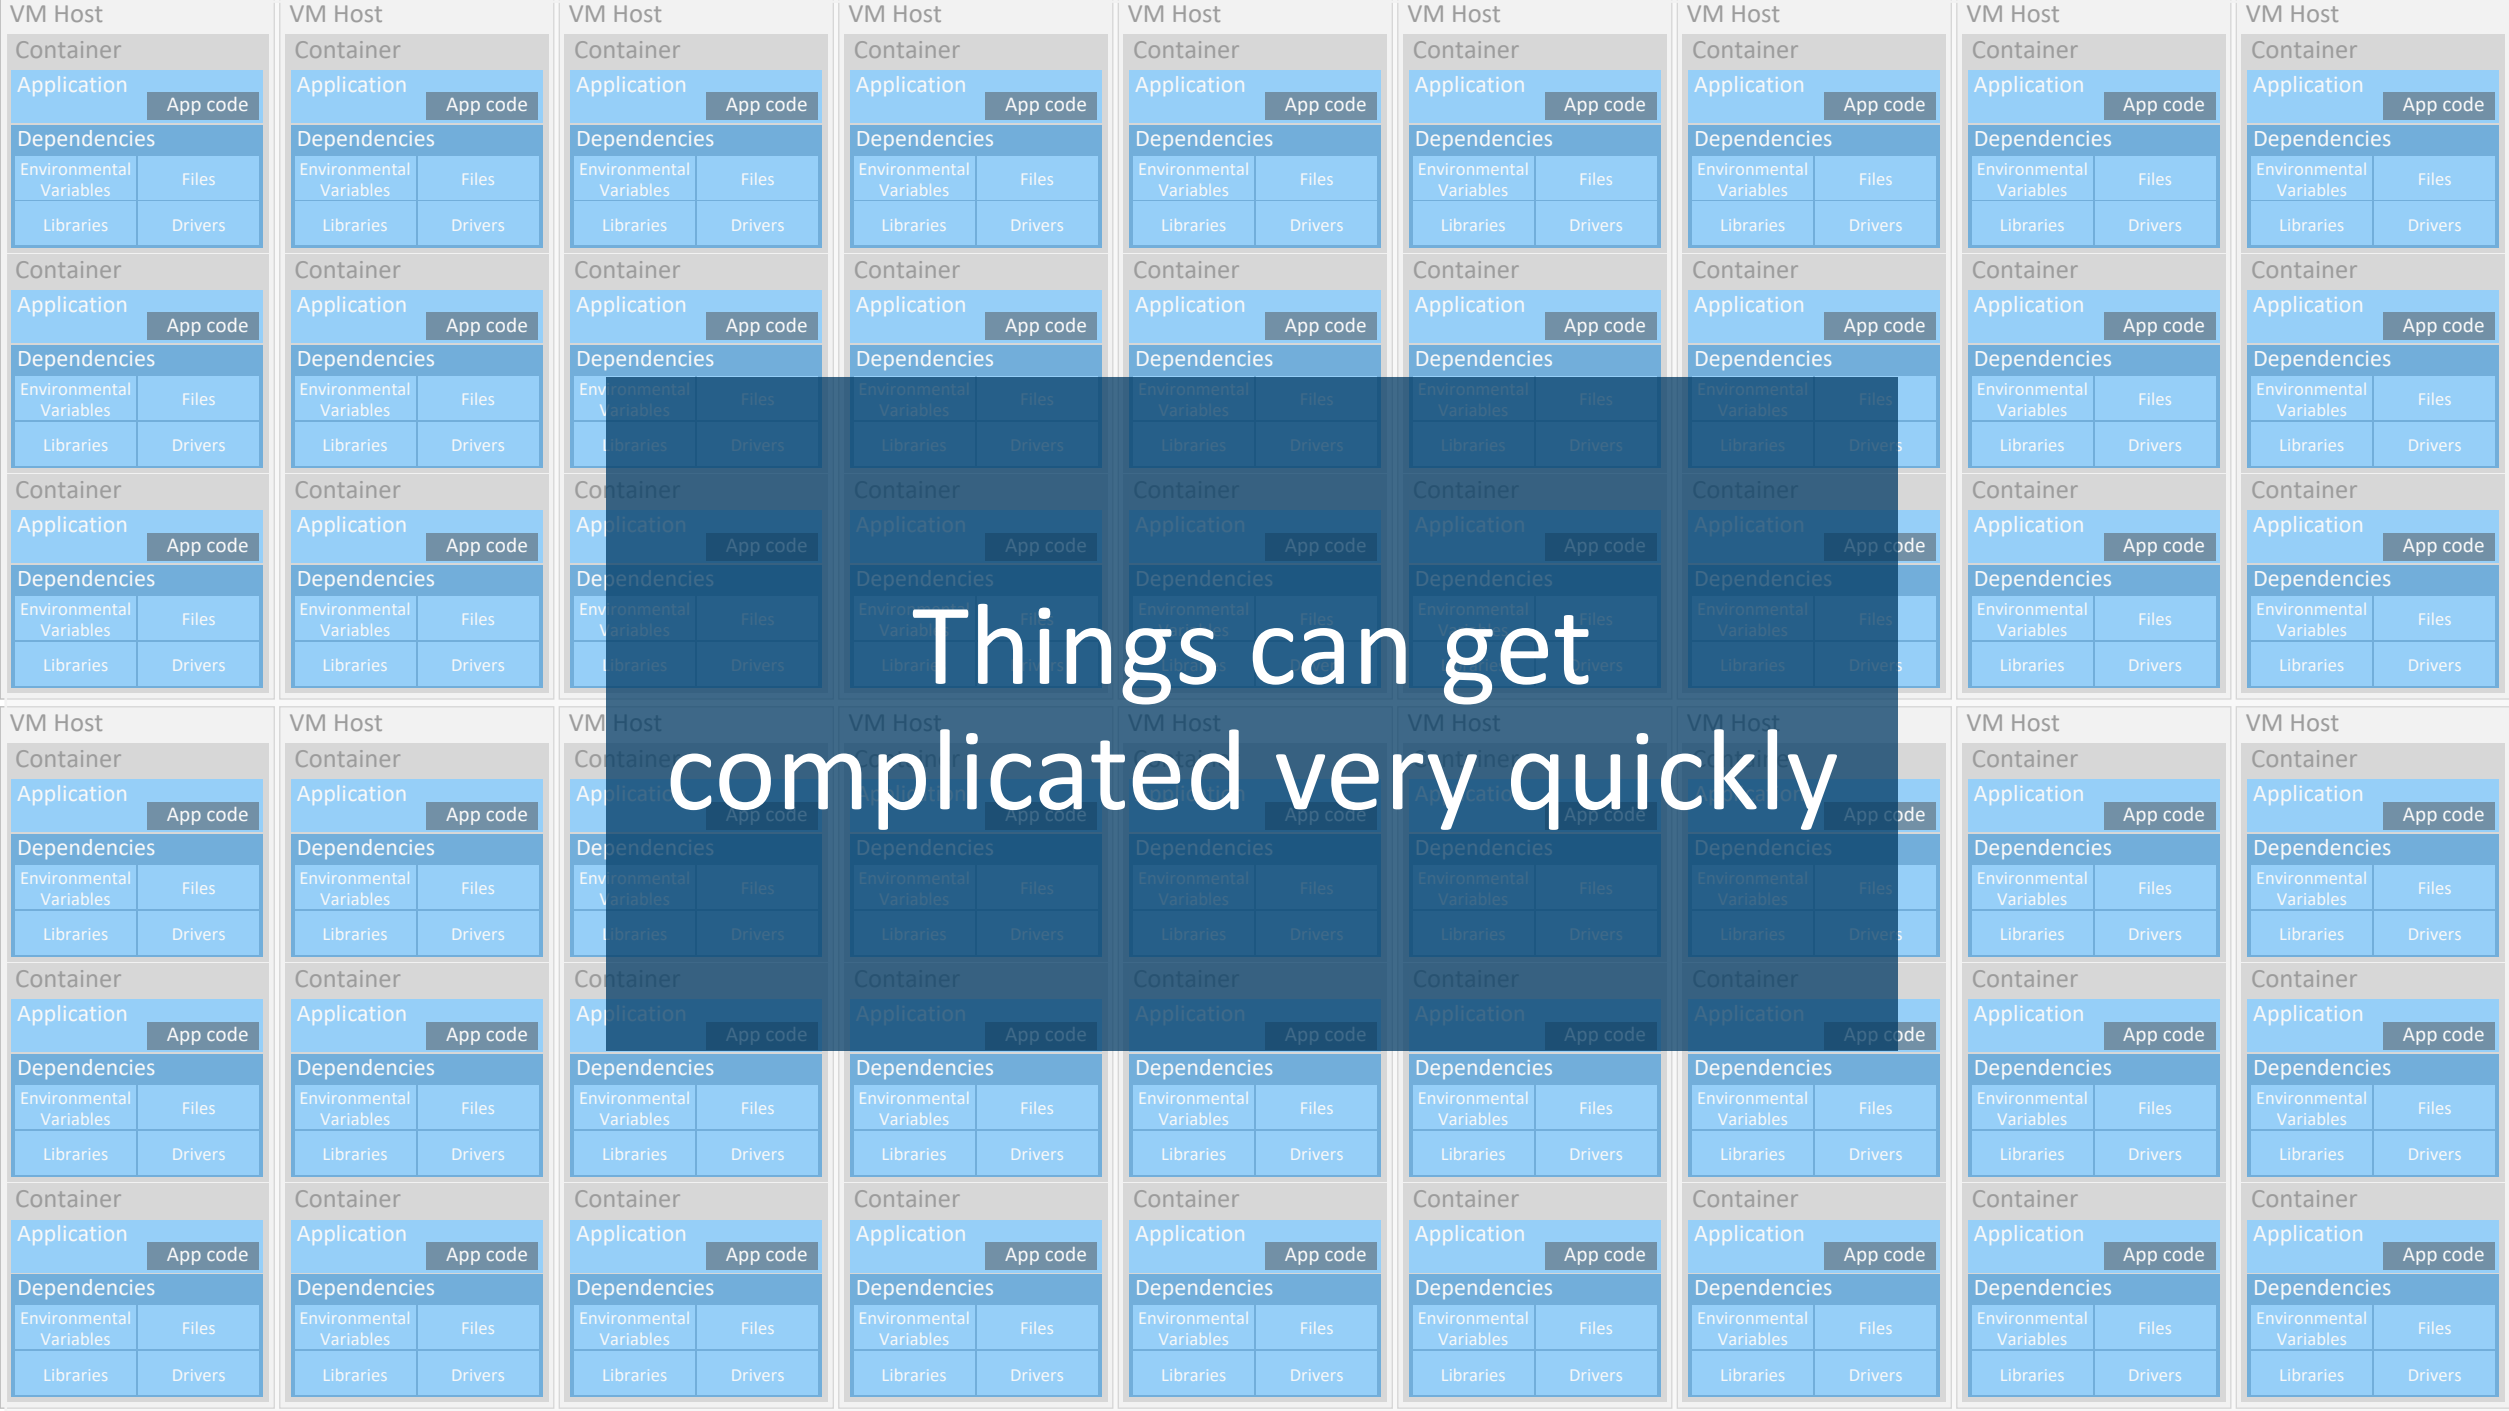
<https://aka.ms/Kubernetes-java>

# Agenda

- Why orchestration is needed?
- What is Kubernetes?
- The history and evolution of Kubernetes
- Kubernetes architecture and components
- Toolset
- Demo

# Container Challenges





# Container Management at Scale

**Cluster Management:**  
deploy and manage cluster resources

**Scheduling:**  
where containers run

**Lifecycle and Health:**  
keep containers running despite failure

**Naming and Discovery:**  
where are my containers

**Load Balancing:**  
evenly distribute traffic

At the end of the day we need something to help us with all the orchestration..  
**An orchestrator!**

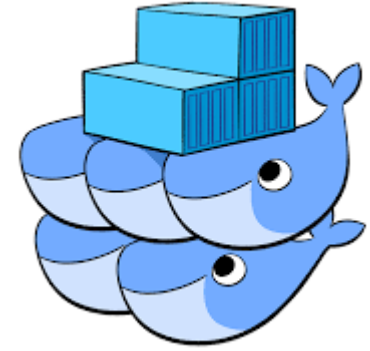
**Scaling:**  
make sense of container in numbers

container images

containers and cluster

# Available Orchestrators

- Docker Swarm
- Apache Mesos
- Nomad (from HashiCorp)
- Rancher
- Service Fabric
- ...
- Kubernetes



MESOS



HashiCorp  
**Nomad**



**RANCHER**



Microsoft Azure  
Service Fabric

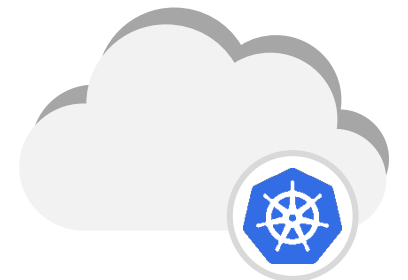


# Kubernetes is Born



# What is Kubernetes (**k8s**)?

- **Kubernetes** is "an open-source software for automating deployment, scaling, and management of containerized applications".
- **Kubernetes**, in Greek κυβερνήτης, means the Helmsman, or pilot of the ship.
- Keeping with the maritime theme of **Docker** containers, **Kubernetes** is the pilot of a ship of containers.





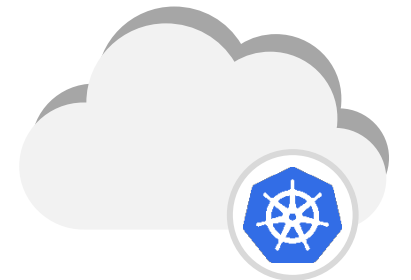
# What is Kubernetes (k8s)?

## History

- Originally designed by Google and is now maintained by the Cloud Native Computing Foundation (CNCF).
- Google still actively involved
- Kubernetes v1.0 was released on July, 2015 by Joe Beda, Brendan Burns and Craig McLuckie
- Most discussed repo in GitHub last year.
- Over 1,700 authors and releases every three month
- To learn more about the ideas behind Kubernetes: read the [Large-scale cluster management at Google with Borg](#) paper

# Kubernetes Features

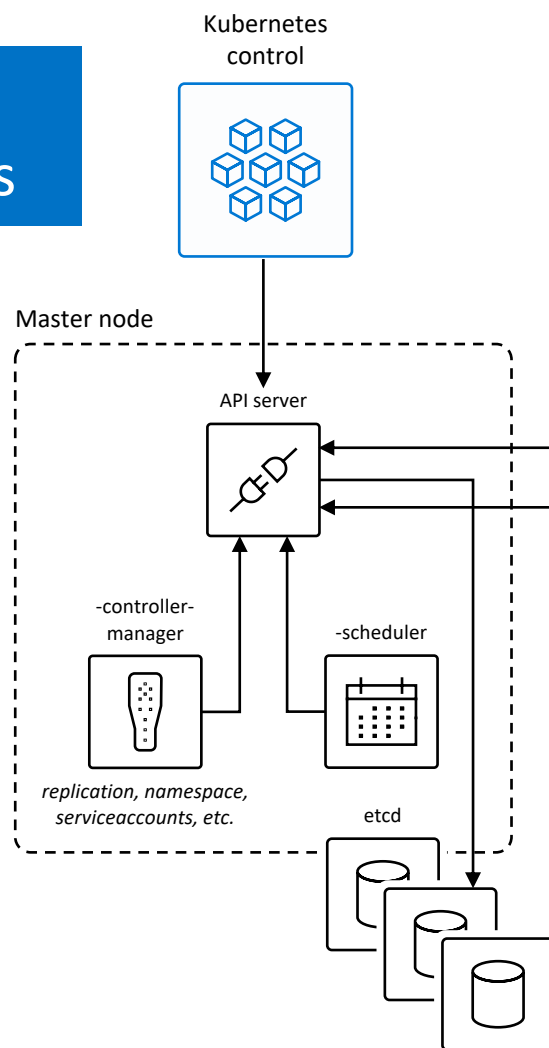
- Self-Healing
- Horizontal Scaling
- Automated rollouts and rollbacks
- Service Discovery and Load Balancing
- Automatic bin packing
- Storage orchestration
- Secret and configuration management



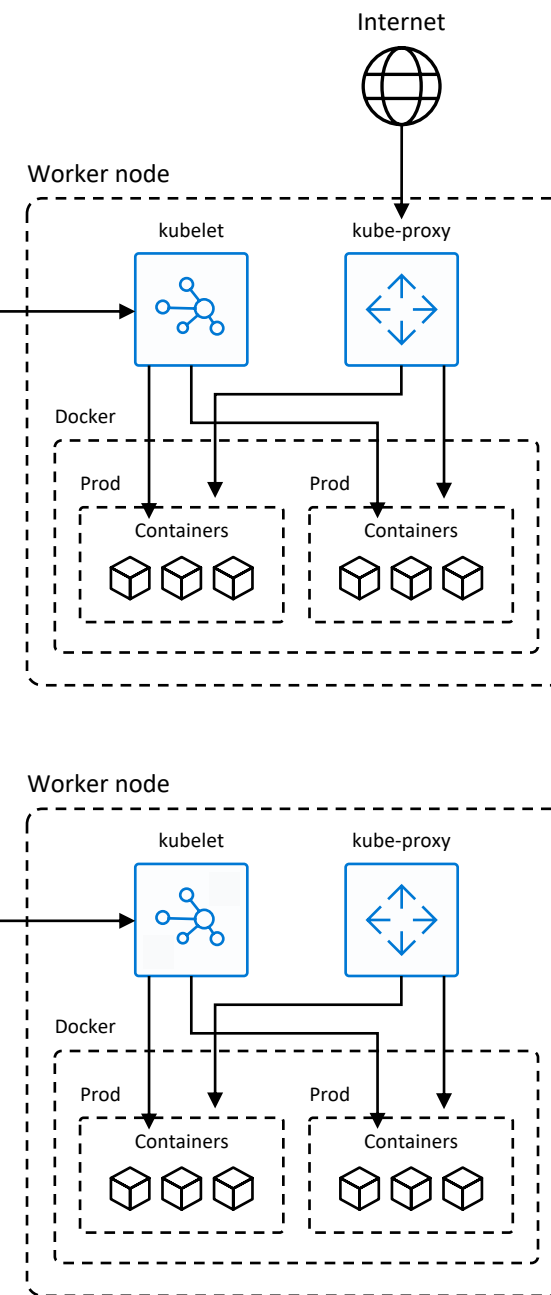
# Kubernetes Building Blocks



## master components



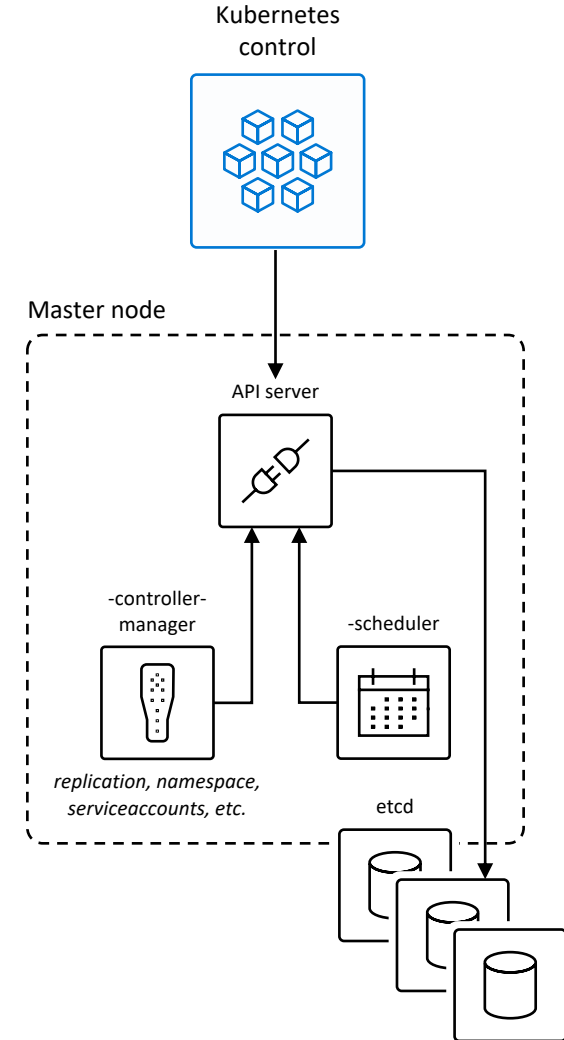
## node components



# kubectl

## master components

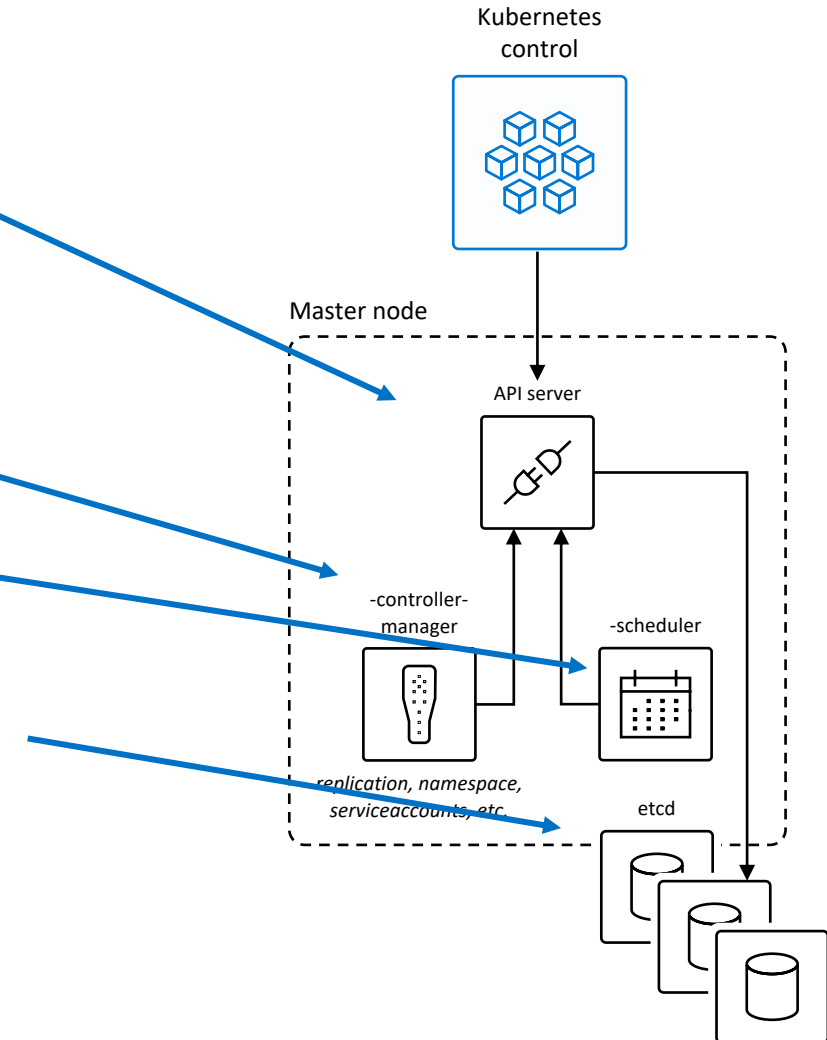
- CLI to run commands against a Kubernetes cluster
- Swiss Army Knife: run deployments, exec into containers, view logs, etc.
- Pronounced “koob sea tee el” or “koob cuddle”
- Available for Windows and Linux – of course available in Azure Cloud Shell



# Master Components

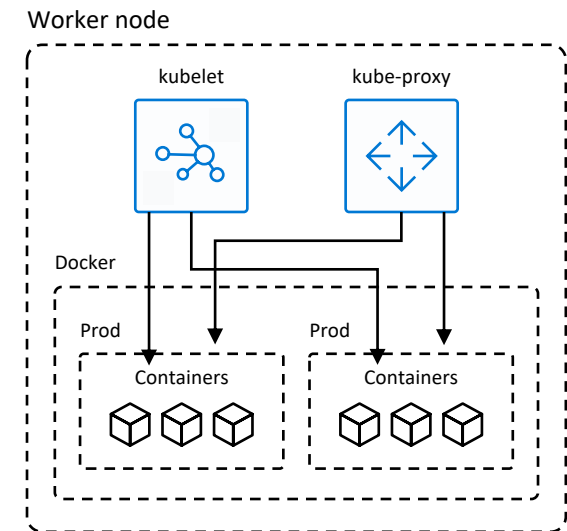
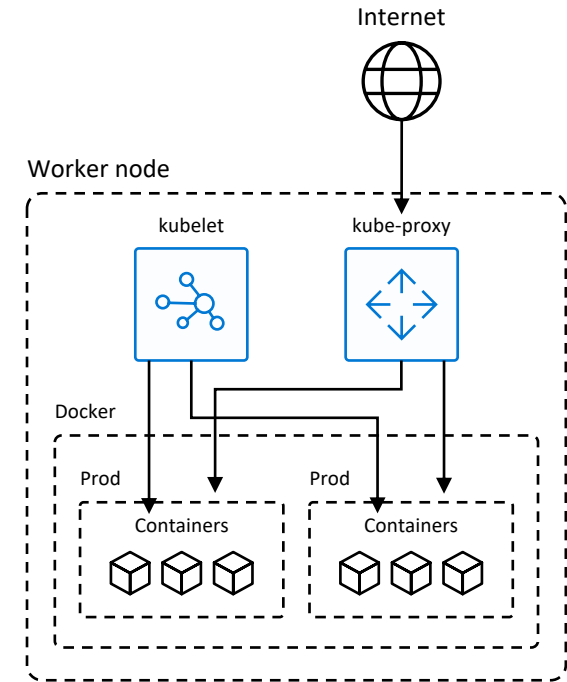
master  
components

- **Kube-api-server**
  - Front-end control plane. Exposes API
- **controller-manager**
  - Runs controllers, e.g. replication controller, node controller
- **scheduler**
  - assigns pods to nodes
- **etcd**
  - Highly available, distributed Cluster database.
- **add-ons**
  - DNS, Heapster (enables monitoring and performance analysis), Dashboard, Logging



# Worker Node Compos

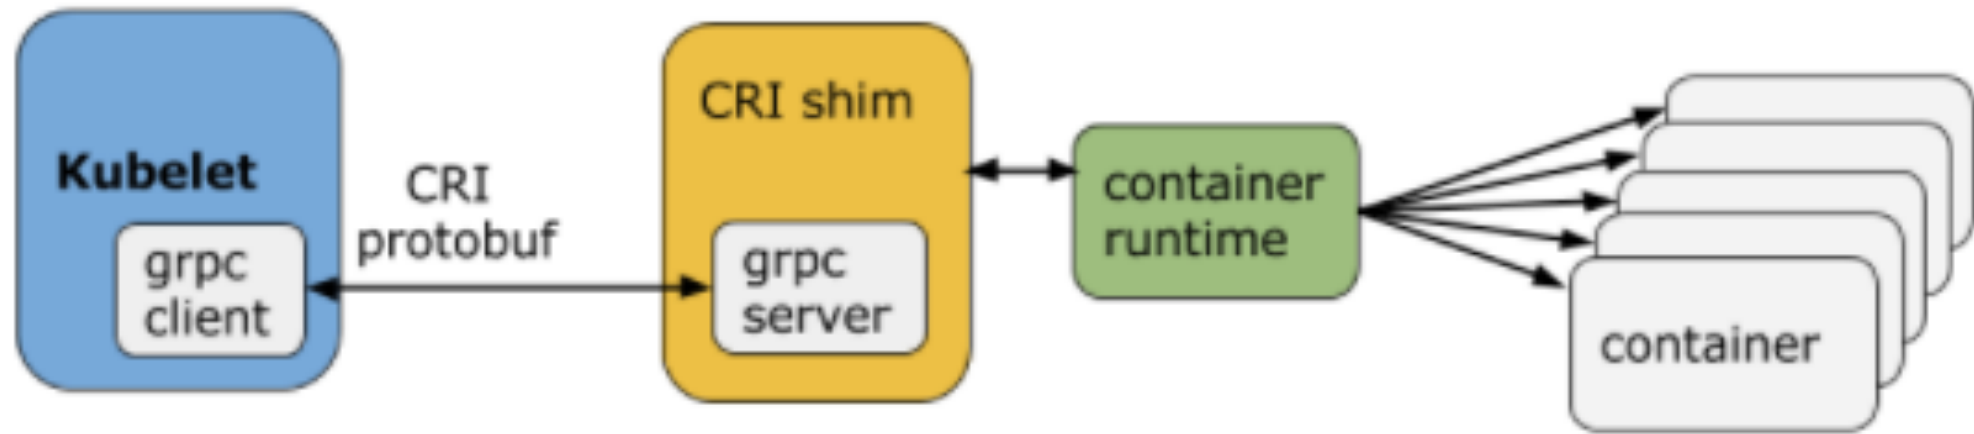
- **Kubelet**
  - Primary node agent
  - Watches and runs assigned pods
  - Executes health probes and reports status
- **Kube-proxy**
  - Enables network services
- **Container Runtime**
  - Docker, rkt, runc and other OCI implementations..



node  
components

# Worker Node Components

- **Container Runtime**
  - Kubernetes has **no code** to execute or run containers on Linux or Windows
  - Initially the Kubernetes pod manager (called “kubelet”) had direct linkage to the **Docker** engine
  - Container Runtime Interface (CRI) was introduced with Kubernetes 1.5





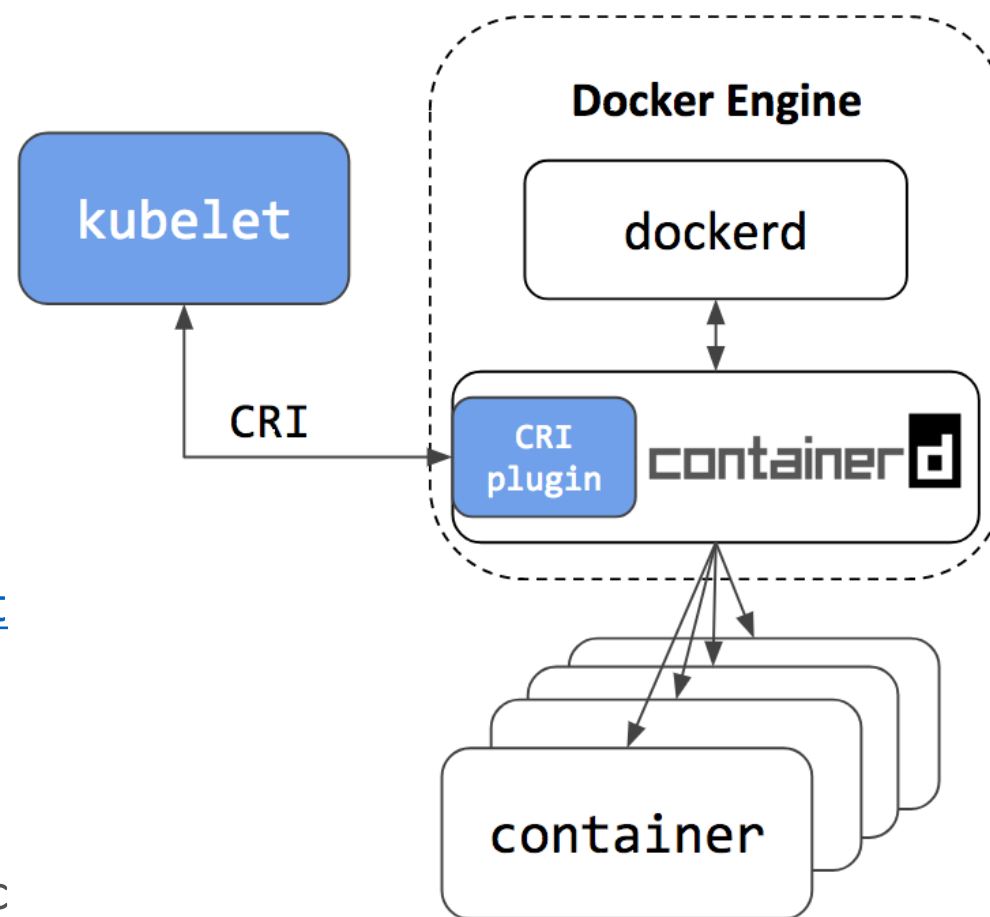
# Worker Node Components

- Container Runtime Interface (CRI)
  - provides a clearly-defined abstraction layer
  - eliminates barriers building own container runtimes
  - enabling pluggable container runtimes

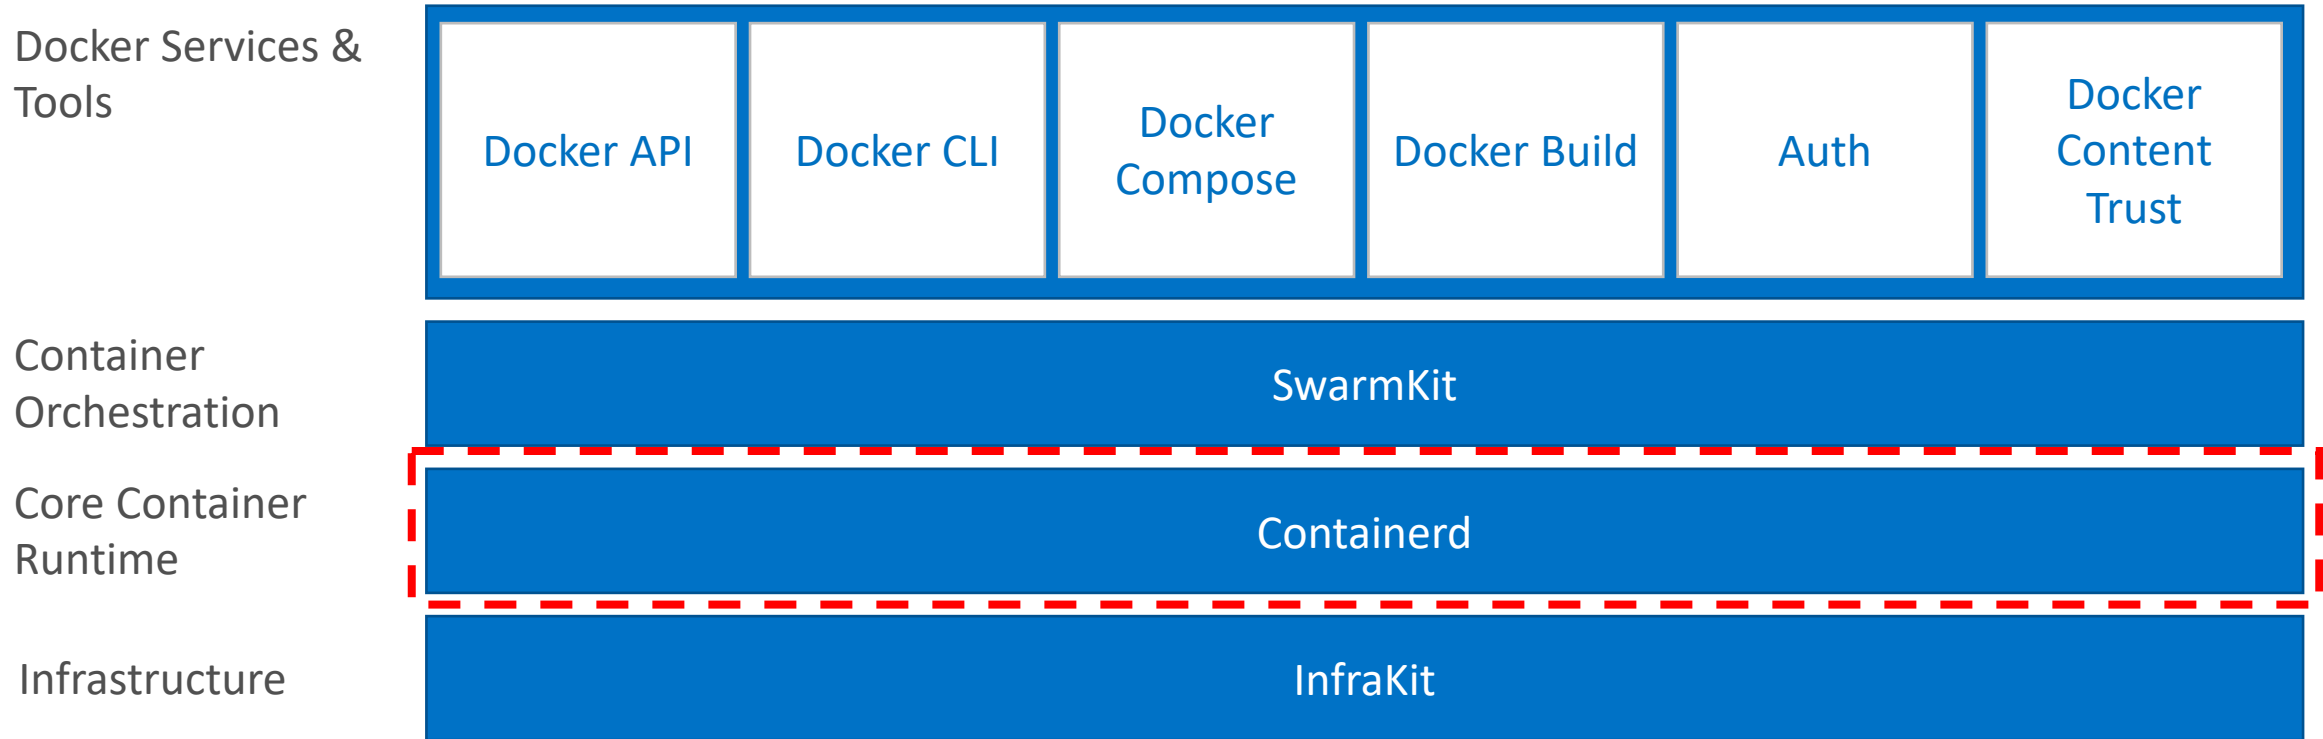
More information about CRI can be found at -

<https://kubernetes.io/blog/2016/12/container-runtime-int-in-kubernetes/>

Runtimes supported today, either upstream or by forks, include docker (for Linux and Windows), [rkt](#), [cri-o](#), and [frakti](#).



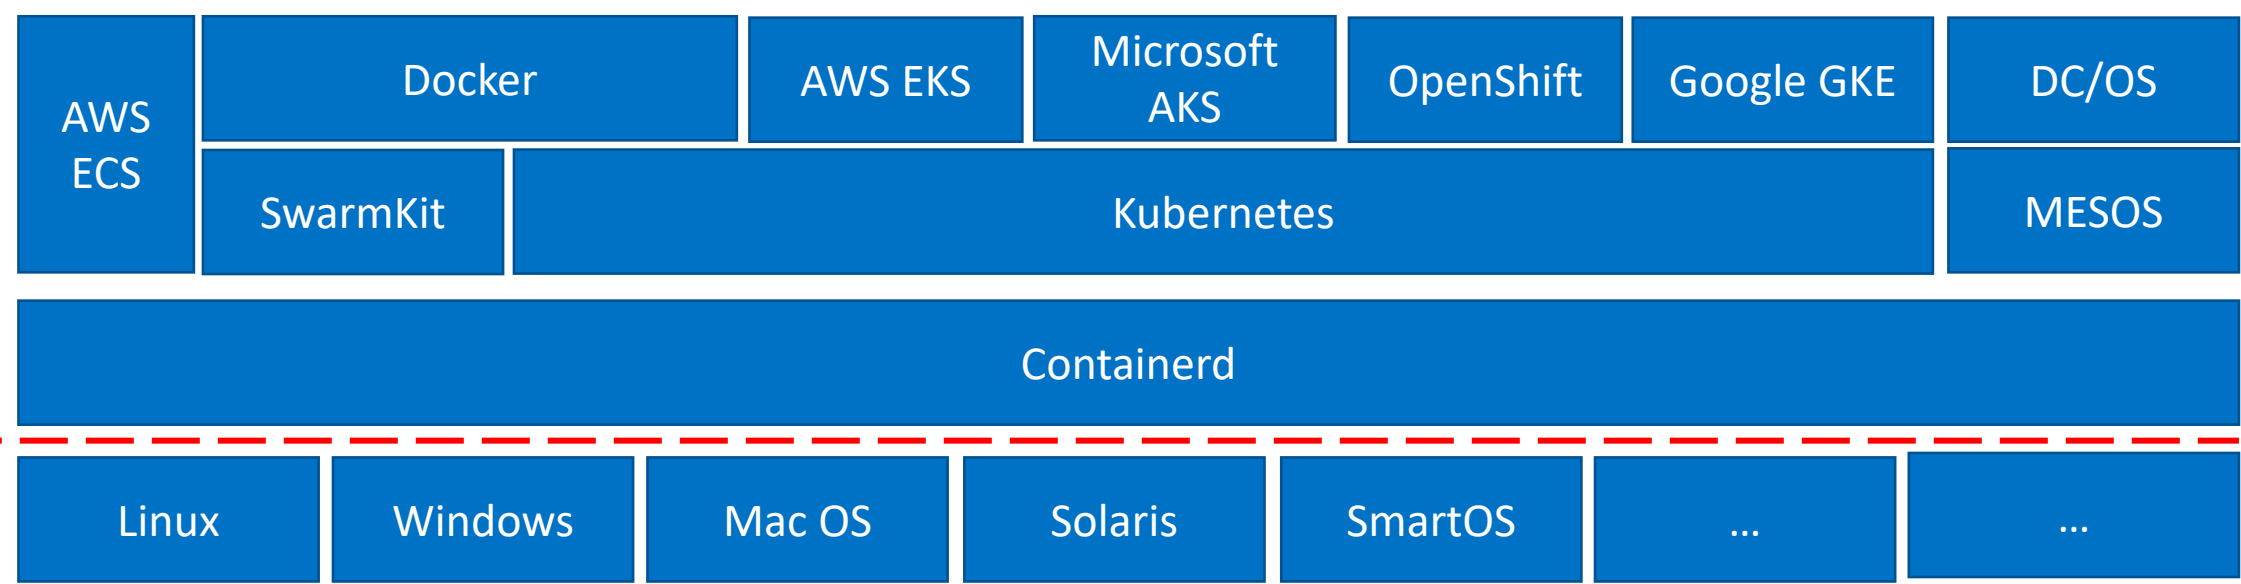
# Docker Engine



DOCKER TO DONATE CONTAINERD TO THE CLOUD NATIVE COMPUTING FOUNDATION (MARCH 2017)

<https://blog.docker.com/2017/03/docker-donates-containerd-to-cncf/>

# Containerd's role in Container Ecosystem



containerd deep dive @ containerd summit for more  
[https://www.youtube.com/watch?v=UUDDCetB7\\_A](https://www.youtube.com/watch?v=UUDDCetB7_A)

\*OCI = Open Container Interface

# What is Moby?



- **Containerd** – core container runtime
- **Linuxkit** – tool to build secure, portable and lean os for containers.
- **Infrakit** – creating and managing self healing infrastructures.
- **Libnetwork** – native Go implementation to connect containers.

## About Moby:

“An open framework to assemble specialized container systems without reinventing the wheel.”

[mobyproject.org](https://mobyproject.org)

## Why do we care?

“To deliver more frequent upstream patches and improvements to the container runtime, AKS has adopted Moby, the open-source project that Docker is based on. “

<https://azure.microsoft.com/en-us/updates/azure-kubernetes-service-december-update/>

# Kubernetes Resources



# Declarative vs. Imperative

- Commands like `kubectl run` and `kubectl expose` are **imperative** commands (do this thing now)

```
$> kubectl run -i --tty busybox --image=busybox --restart=Never -- sh
```

- **Declarative** way – Describe the state of resources in a file (JSON or YAML).

`kubectl apply -f webresource.yaml`

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
```

# Kubernetes Resources

pod

ReplicaSet

deployment

service

namespace

volumes

config-map

secret

ingress

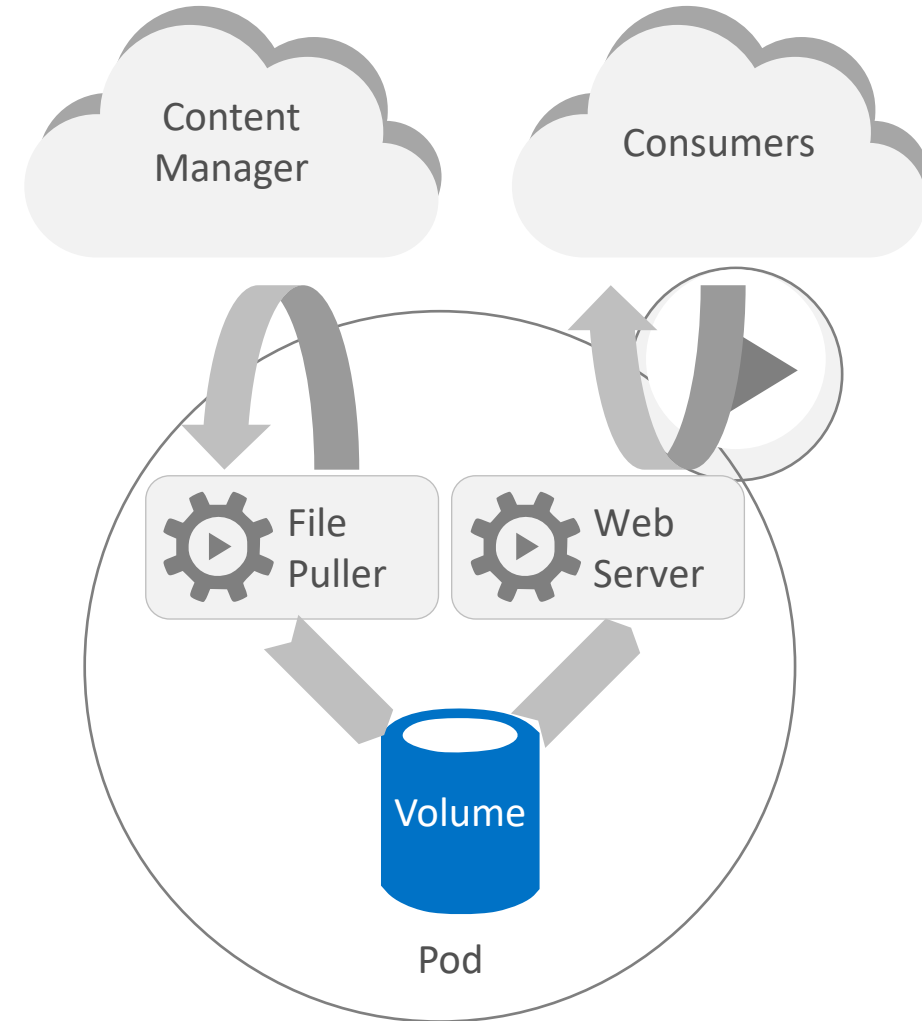
StatefulSet

DaemonSet

jobs

# What is a pod?

- Pod is the basic building block in Kubernetes
- Pods are how containers are delivered
- Can be multiple containers (e.g. side car)
- Encapsulates container(s), storage, network IP, and options on how to run





# What is a pod? Considerations..

- Pods do not, by themselves, self-heal!
  - If a Pod is scheduled to a Node that fails, or if the scheduling operation itself fails, the Pod is deleted; likewise, a Pod won't survive an eviction due to a lack of resources or Node maintenance.
- Thus, while it is possible to use Pod directly, it's far more common in Kubernetes to manage your pods using a Controller.
- Controllers can create and manage multiple Pods, handling replication and rollout and providing self-healing capabilities at cluster scope.
  - For example, if a Node fails, the Controller might automatically replace the Pod by scheduling an identical replacement on a different Node.

# What is a pod?

- Examples of Controllers that contain one or more pods:
  - StatefulSet
  - DaemonSet
  - ReplicaSet (+ Deployments)
  - Jobs

deployment

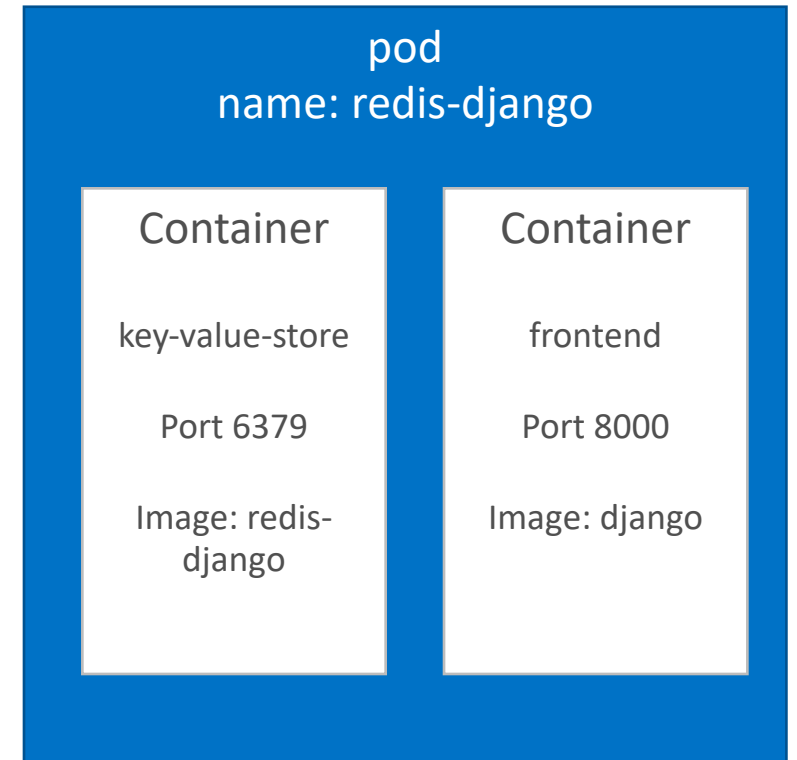
StatefulSet

DaemonSet

ReplicaSet

# Kubernetes manifest: Pod

```
apiVersion: v1
kind: Pod
metadata:
  name: redis-django
  labels:
    app: web
spec:
  containers:
    - name: key-value-store
      image: redis
      ports:
        - containerPort: 6379
    - name: frontend
      image: django
      ports:
        - containerPort: 8000
```



# Kubernetes manifest: Pod

```
apiVersion: v1
kind: Pod
metadata:
  name: redis-django
  labels:
    app: web
spec:
  containers:
  - name: key-value-store
    image: redis
    resources:
      requests:
        memory: "64Mi"
        cpu: "250m"
      limits:
        memory: "128Mi"
        cpu: "500m"
```

- **Requests** specify how much resources a container needs.
- **Limits** define how much resources a container can consume.

# Interact with pods

```
$ kubectl get pod --all-namespaces
```

```
$ kubectl describe pod/my-pod
```

```
$ kubectl logs my-pod
```

```
# Run ba
```

```
$ kubect
```

```
heyko@Azure:~$ kubectl describe pod/brigade-brigade-api-65b74b4cc8-6f8ws
Name:          brigade-brigade-api-65b74b4cc8-6f8ws
Namespace:     default
Node:          aks-nodepool11-60876065-0/10.240.0.5
Start Time:    Mon, 09 Jul 2018 13:12:48 +0000
Labels:        app=brigade-brigade-api
```

```
heyko@Azure:~$ kubectl logs brigade-brigade-api-65b74b4cc8-6f8ws
```

```
[GIN-debug] [WARNING] Running in "debug" mode. Switch to "release" mode in production.
```

```
- using env:  export GIN_MODE=release
```

```
- using code: gin.SetMode(gin.ReleaseMode)
```

```
[GIN-debug] GET    /v1/projects          --> github.com/Azure/brigade/pkg/api.(Project).List-fm (4 handlers)
```

```
[GIN-debug] GET    /v1/project/:id       --> github.com/Azure/brigade/pkg/api.(Project).Get-fm (4 handlers)
```

```
[GIN-debug] GET    /v1/project/:id/builds --> github.com/Azure/brigade/pkg/api.(Project).Builds-fm (4 handlers)
```

```
[GIN-debug] GET    /v1/projects-build    --> github.com/Azure/brigade/pkg/api.(Project).ListWithLatestBuild-fm
```

```
[GIN-debug] GET    /v1/build/:id         --> github.com/Azure/brigade/pkg/api.(Build).Get-fm (4 handlers) ce2663b19efe860aa679d
```

```
[GIN-debug] GET    /v1/build/:id/jobs    --> github.com/Azure/brigade/pkg/api.(Build).Jobs-fm (4 handlers)
```

```
[GIN-debug] GET    /v1/build/:id/logs    --> github.com/Azure/brigade/pkg/api.(Build).Logs-fm (4 handlers)
```

```
[GIN-debug] GET    /v1/job/:id           --> github.com/Azure/brigade/pkg/api.(Job).Get-fm (4 handlers)
```

```
[GIN-debug] GET    /v1/job/:id/logs      --> github.com/Azure/brigade/pkg/api.(Job).Logs-fm (4 handlers)
```

```
[GIN-debug] GET    /healthz              --> github.com/Azure/brigade/pkg/api.Healthz (2 handlers)
```

```
[GIN-debug] Listening and serving HTTP on :7745
```

```
BRIGADE_NAMESPACE: default (v1:metadata.namespace)
```

```
BRIGADE_API_PORT: 7745
```

```
Mounts:
```

```
/var/run/secrets/kubernetes.io/serviceaccount from brigade-brigade-api-token-px7qr (ro)
```

# Labels and Selectors

- Labels are key/value pairs for any API object in Kubernetes
- "Label selectors" == queries against labels to match objects
- Use cases:
  - Associating pods to a service
  - Pinning workloads to specific nodes
  - Selecting a subset of resources

```
apiVersion: v1
kind: Pod
metadata:
  name: example-pod
  labels:
    env: development
spec:
  containers:
  - name: label-example
    image: sonasingh46/node-web-app:latest
    ports:
    - containerPort: 8000
```

# Controllers - Deployment

Deployment provides declarative updates for Pods and ReplicaSets.

Use Cases:

- Create deployment to rollout **ReplicaSet**
- Declare new state for pods (eg – new imageTag)
- Rollback to earlier revision
- Scale up or down
- Check rollout history
- Clean-up older ReplicaSets

# Kubernetes manifest: Deployment

```
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: smackweb-deploy
spec:
  replicas: 4
  template:
    metadata:
      labels:
        app: smackweb
    spec:
      containers:
      - name: smackweb
        image: chzbrgr71/smackweb
        ports:
        - containerPort: 8080
```

pod  
name: smackweb-  
deploy-xyz1

Container

smackweb  
Port 8080

Image: smackweb  
labels: smackweb

pod  
name: smackweb-  
deploy-xyz2

Container

smackweb  
Port 8080

Image: smackweb  
labels: smackweb

pod  
name: smackweb-  
deploy-xyz3

Container

smackweb  
Port 8080

Image: smackweb  
labels: smackweb

pod  
name: smackweb-  
deploy-xyz4

Container

smackweb  
Port 8080

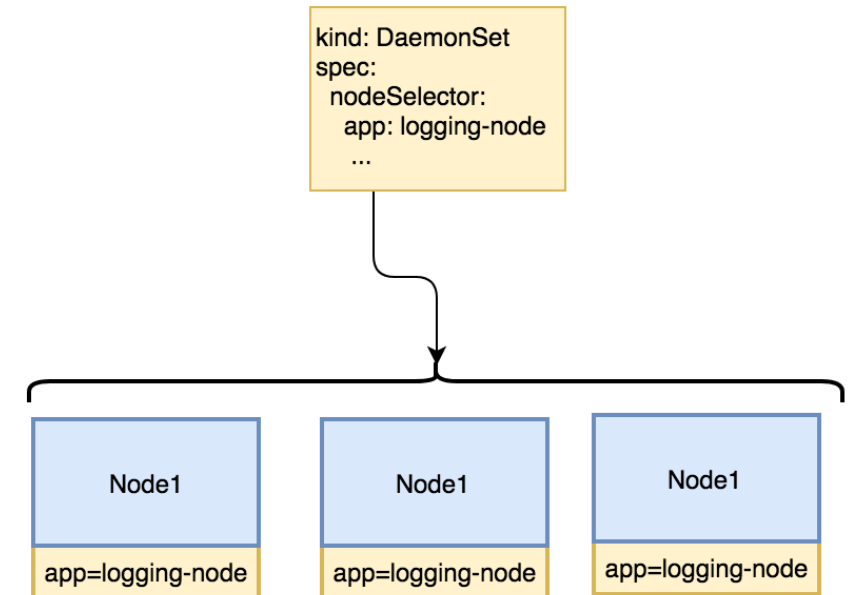
Image: smackweb  
labels: smackweb



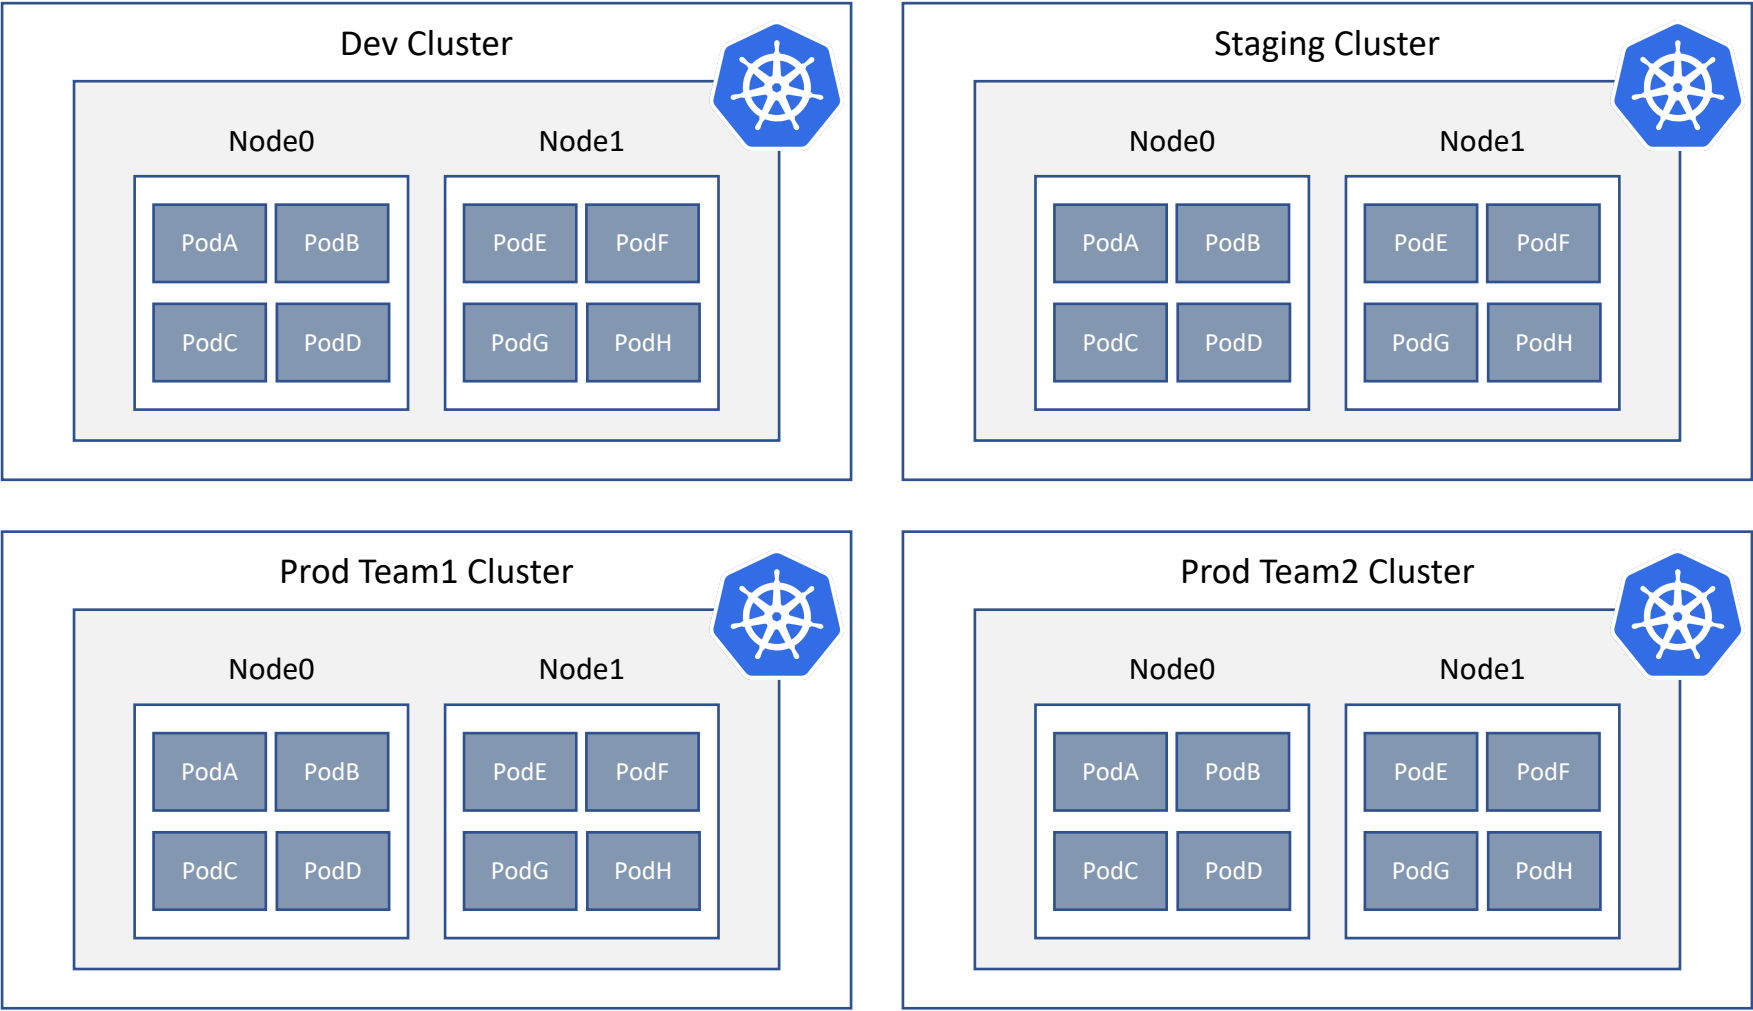
# Controllers - DaemonSets

DaemonSets ensure that all (or some) Nodes run a copy of a Pod.

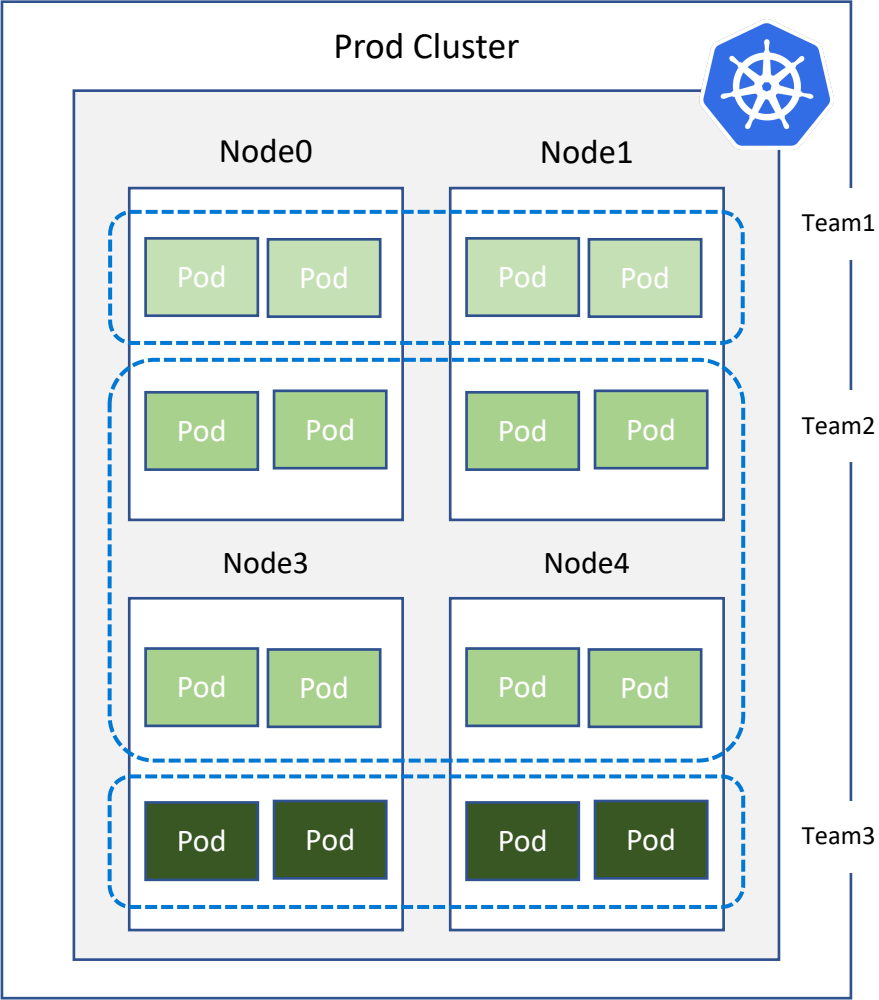
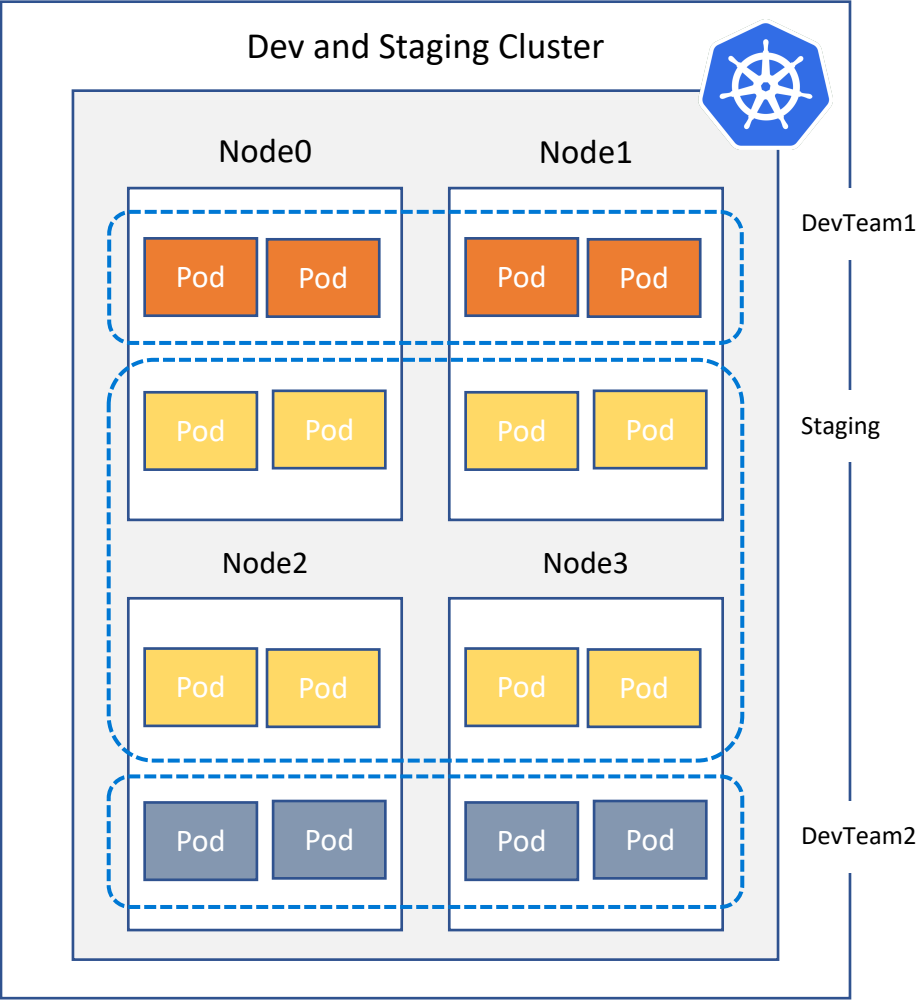
- As worker nodes are
  - added to the cluster, Pods are added to them.
  - removed from the cluster, those Pods are garbage collected.
- Some typical uses of a DaemonSet are:
  - logs collection daemon (i.e. fluentd, logstash)
  - Malware scan (install AV)
  - node monitoring daemon (i.e. Prometheus, collectd, Datadog, New Relic)



# Cluster Isolation Patterns: Physical Isolation



# Cluster Isolation Patterns: Logical Isolation



# Namespaces

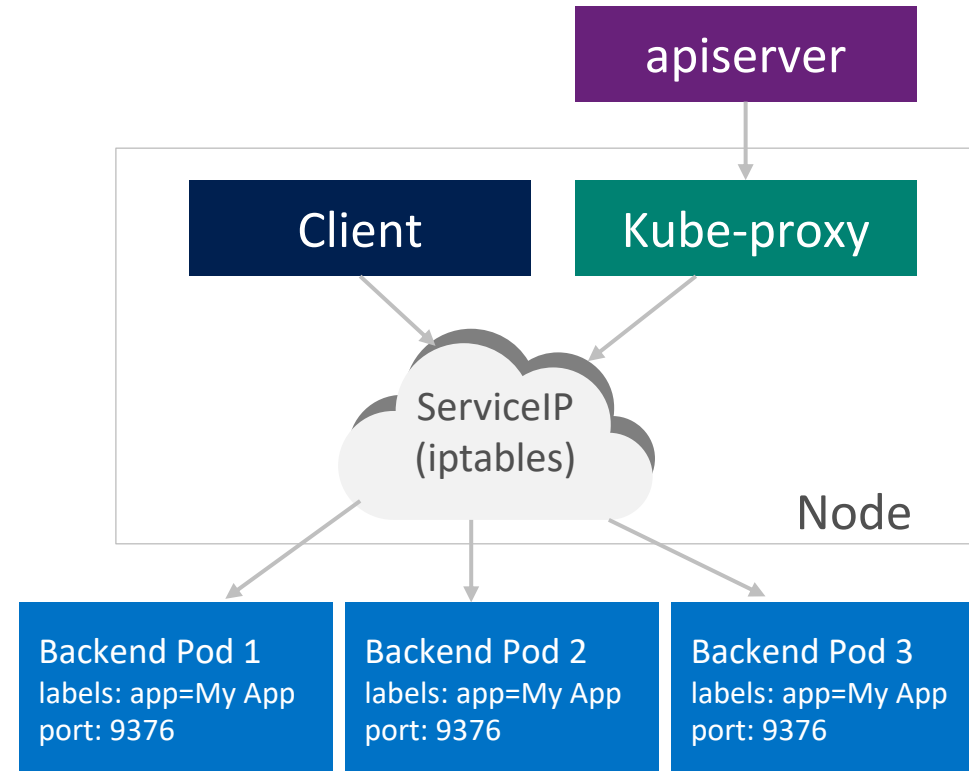
- multiple virtual clusters backed by the same physical cluster
- logical separation/isolation
- Every resource type is scoped to a namespace (except for nodes, persistentVolumes, etc.)
- Intended for environments with many users, teams, projects
- Kube-system namespace for dashboard etc.

```
wslroot@MININT-084LOJC:~$ kubectl get namespaces
```

NAME	STATUS	AGE
default	Active	3d
kube-public	Active	3d
kube-system	Active	3d

# Kubernetes Services

- Defines a logical set of pods
- Identified/selected using Labels

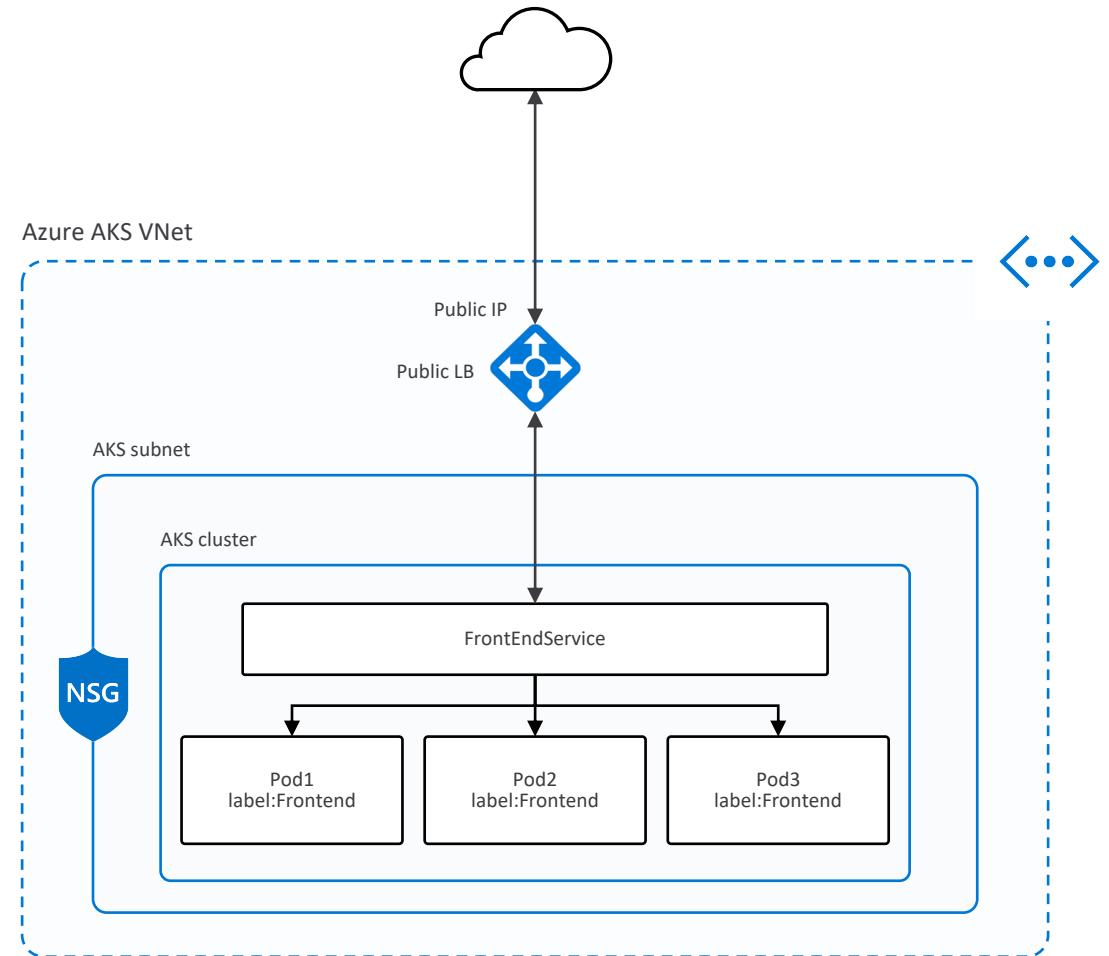


- Essentially a virtual load balancer in front of pods

# Public LoadBalancer Service

- Service Type LoadBalancer
- Basic Layer4 Load Balancing (TCP/UDP)
- Each service has assigned an IP on the ALB

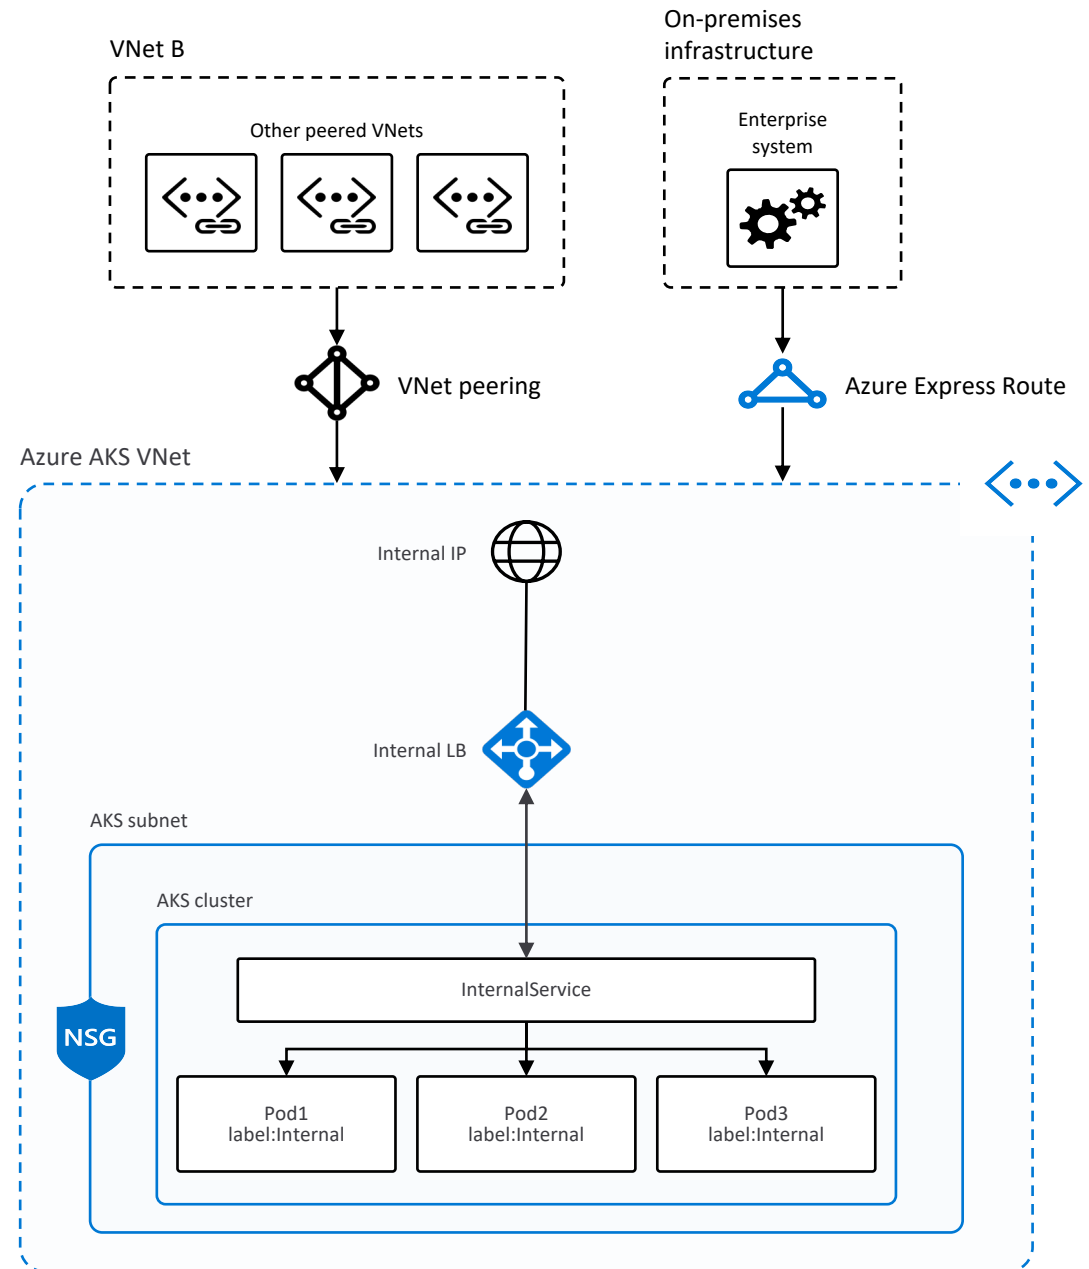
```
apiVersion: v1
kind: Service
metadata:
  name: frontend-service
spec:
  loadBalancerIP: X.X.X.X
  type: LoadBalancer
  ports:
    - port: 80
  selector:
    app: frontend
```



# Internal LoadBalancer Service

- Used for internal services that should be accessed by other VNETs or On-Premise only

```
apiVersion: v1
kind: Service
metadata:
  name: internalservice
  annotations:
    service.beta.kubernetes.io/azure-load-balancer-internal: "true"
spec:
  type: LoadBalancer
  loadBalancerIP: 10.240.0.25
  ports:
    - port: 80
  selector:
    app: internal
```



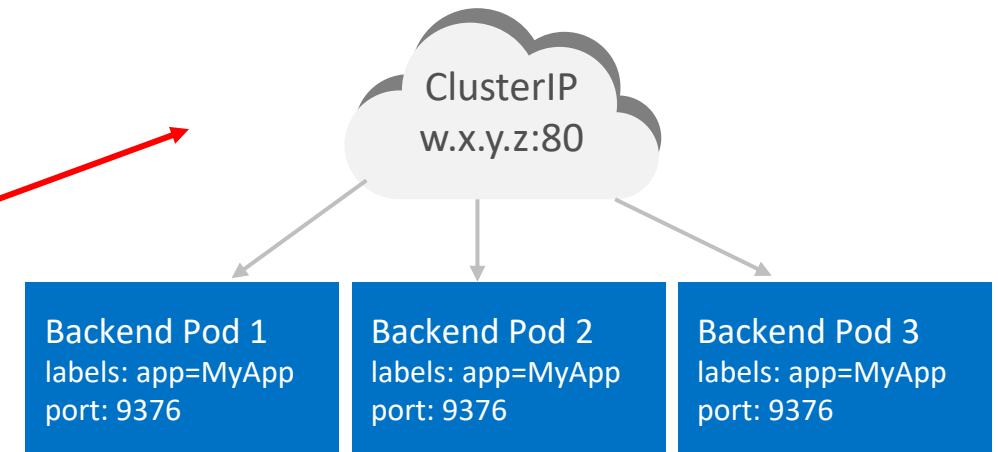
# Other Service Types

- **ClusterIP**
  - Exposes the service on a cluster-internal IP. Choosing this value makes the service **only reachable from within the cluster**
- **NodePort**
  - Exposes the service on each Node's IP at a static port (the NodePort)
  - Connect from outside the cluster by requesting <NodeIP>:<NodePort>



# Kubernetes manifest: Service

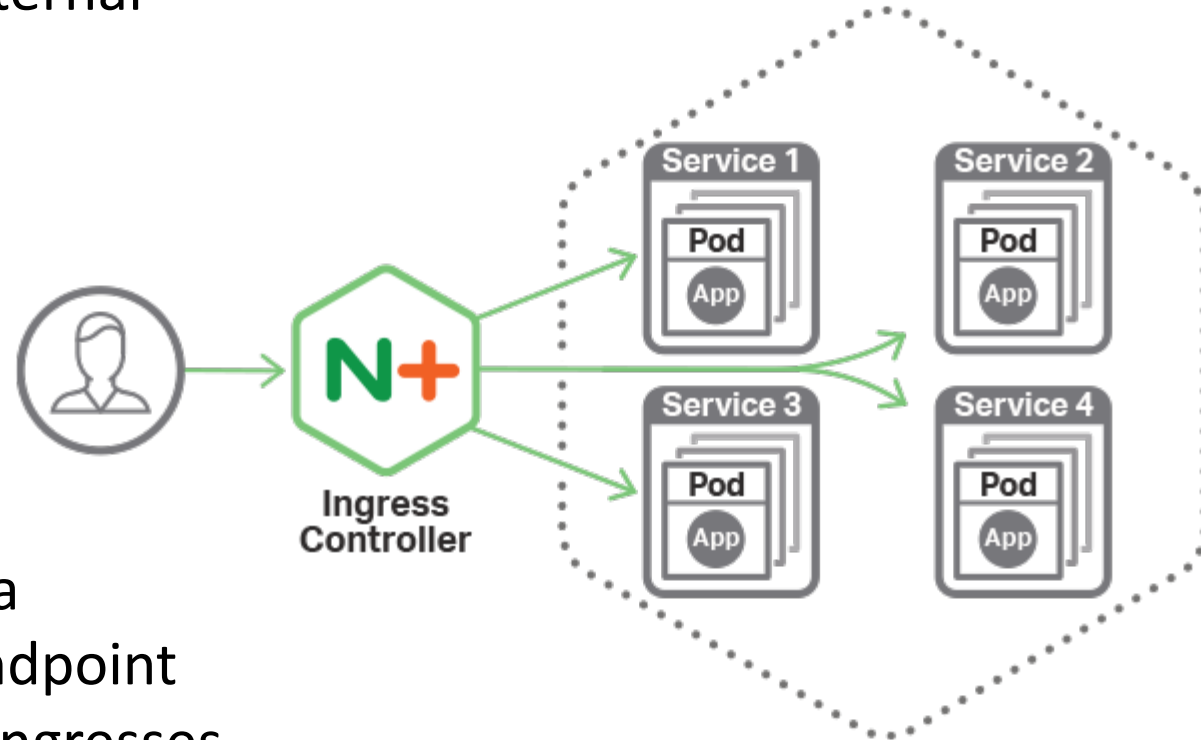
```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: MyApp
  type: ClusterIP
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```



**Note!** Services using **ClusterIP** are only reachable from within the cluster.

# Ingress and Ingress Controllers

- **Ingress** is a Kubernetes API that manages external access to the services in the cluster
  - Supports HTTP and HTTPS
  - Path and Subdomain based routing
  - SSL Termination
  - Save on public Ips
- **Ingress controller** is a daemon, deployed as a Kubernetes Pod, that watches the Ingress Endpoint for updates. Its job is to satisfy requests for ingresses. Most popular one being **Nginx**.



# Secrets, Config Maps

- **Secrets** are intended to hold sensitive information such as passwords, tokens. Secrets are for Confidential data. Secrets are encoded with Base64 encoding
- **ConfigMaps** help you to store non-confidential application configuration data. This helps to decouple configuration artifacts from image content

```
$ kubectl create secret generic db-user-pass --from-file=./username.txt --from-file=./password.txt  
secret "db-user-pass" created
```

# Kubernetes Volumes

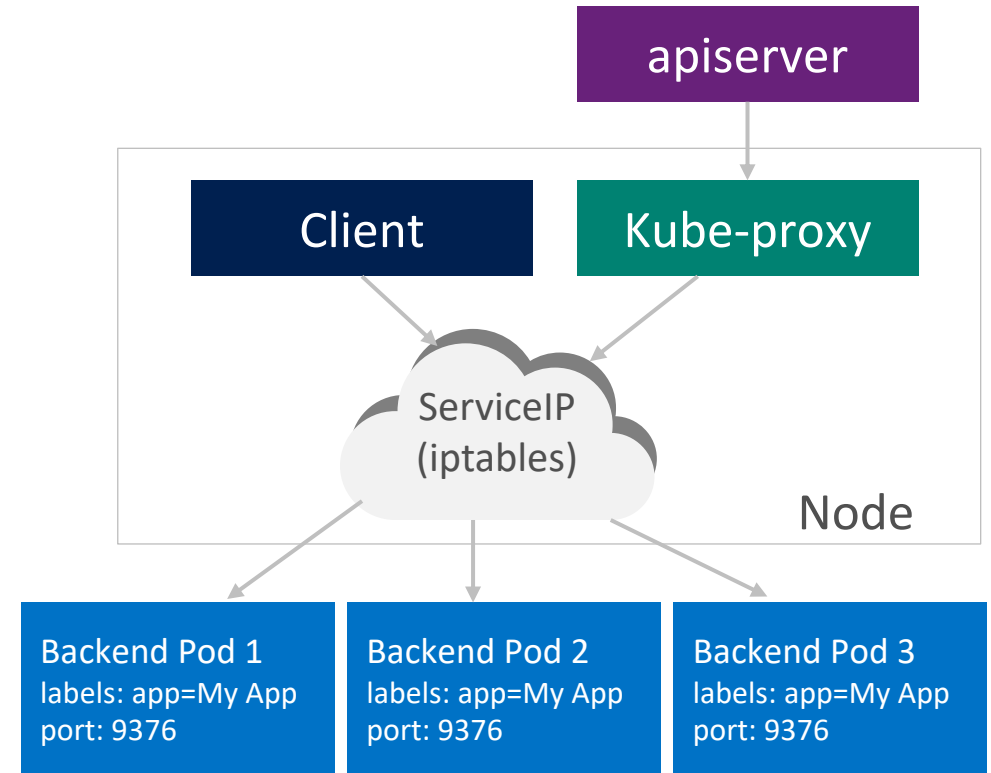
- On disk files in a Container are ephemeral
- Files will be lost if Container crashes and then restarts
- Volumes outlive containers. Lifetime is same as that of a pod. Data is preserved across container restarts
- Persistent Volumes have lifetime independent of the Pod lifetime
- Types of volumes – emptydir, azureDisk, azureFile etc.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: azurefile
provisioner: kubernetes.io/azure-file
mountOptions:
  - dir_mode=0777
  - file_mode=0777
  - uid=1000
  - gid=1000
parameters:
  skuName: Standard_LRS
  storageAccount: mystorageaccount
```

# Networking in Kubernetes

Kubernetes knows 3 methods of communications:

- **Pod-to-Pod** communication directly by IP address. Kubernetes has a Pod-IP wide metric simplifying communication.
- **Pod-to-Service** Communication – Client traffic is directed to service virtual IP by kube-proxy process (running on all hosts) and directed to the correct Pod.
- **External-to-Internal** Communication – external access is captured by an external load balancer which targets nodes in a cluster. The Pod-to-Service flow stays the same.



Demos,  
Labs and  
more



<https://aka.ms/Kubernetes-java>