


# Practical Machine Learning in an Enterprise Setting

Seth Juarez

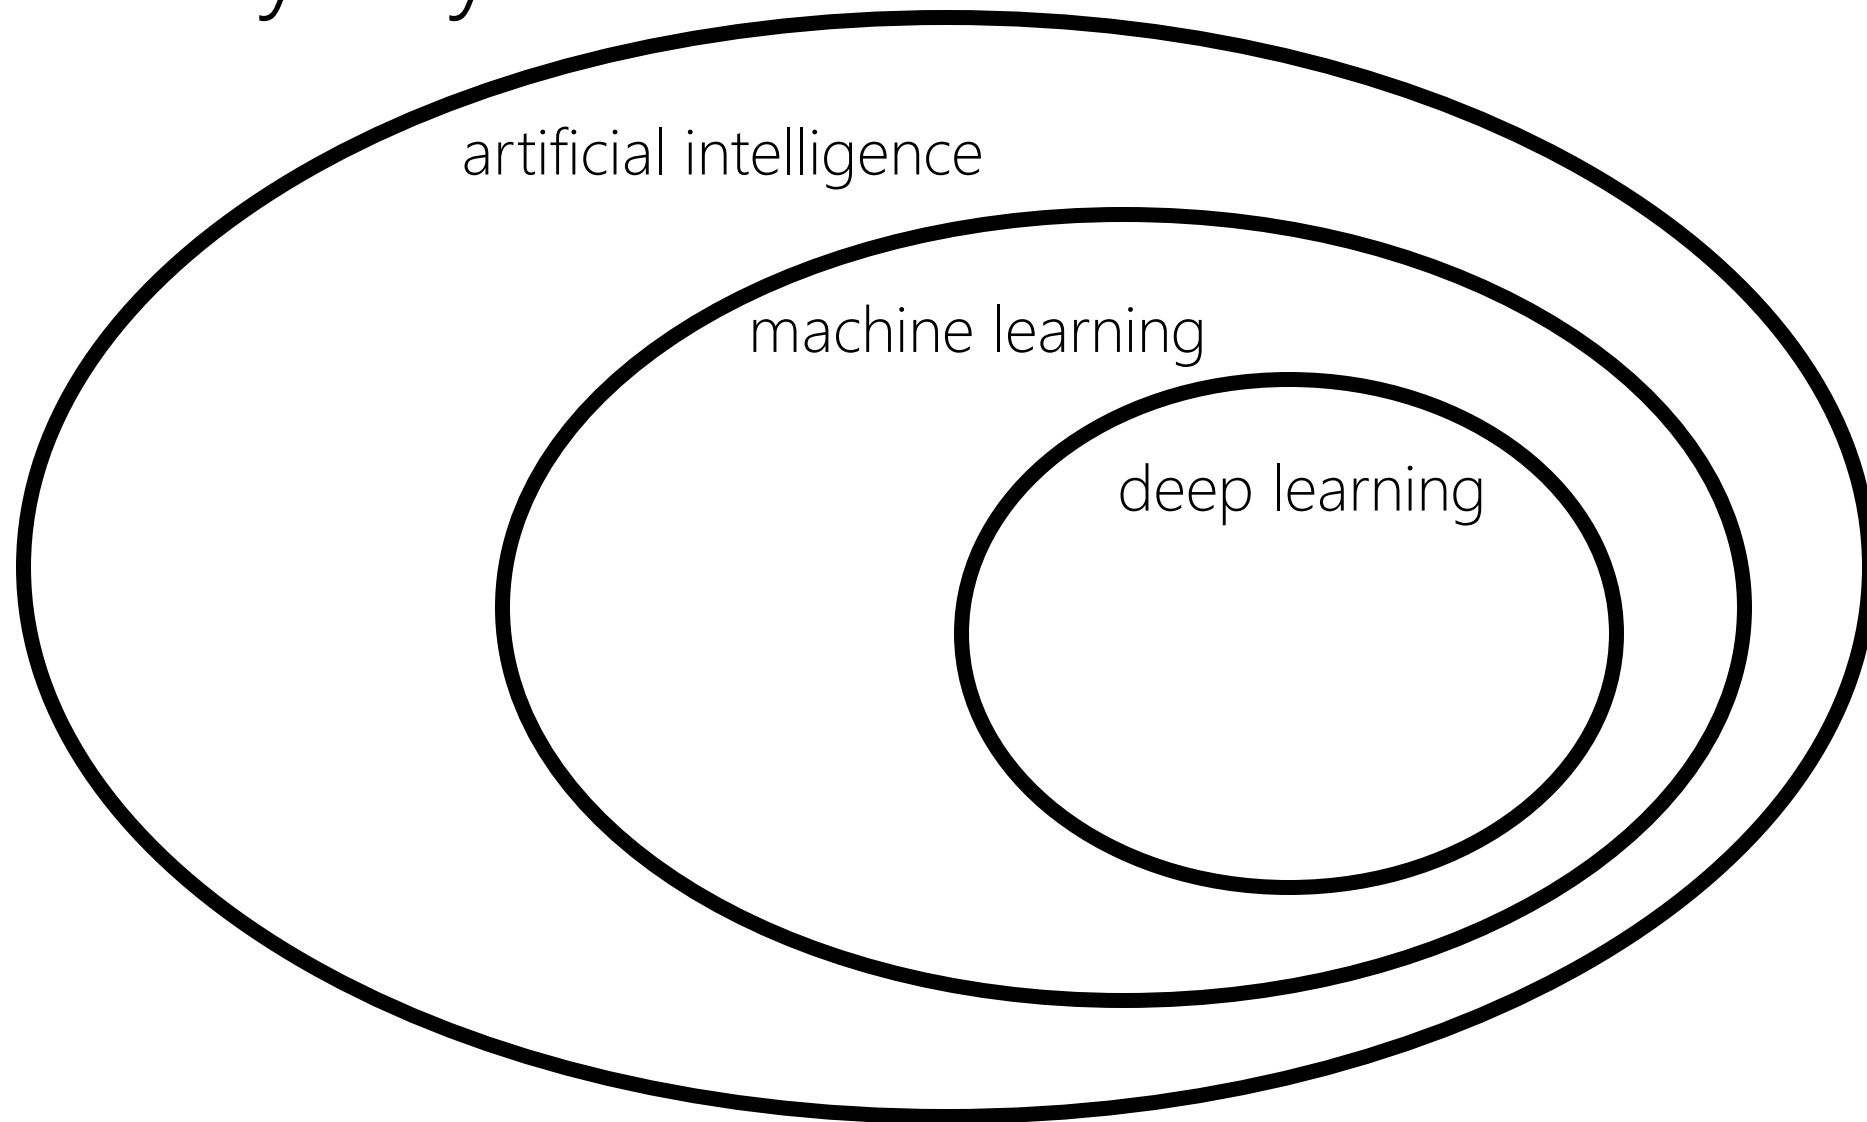
@sethjuarez

# Agenda


- Introduction to Machine Learning
- Three-pronged strategy
- Cloud based machine learning
- AI as a Team Sport – managing the AI Dev Lifecycle




what is it anyway?




when should I use it?




# programming



# machine learning



# machine learning





$h(x)$

scikit-learn

# scikit-learn – concepts


- You “pick” the function “shape” – in scikit-learn these are called *classifiers*

[illegible]

- You “fit” the data to the classifier




- You use the classifier to predict



# scikit-learn – concepts

```
mnist_file = 'https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz'  
file, output = urlretrieve(mnist_file, 'mnist.npz')
```

```
with np.load(file) as f:  
    x_train, y_train = f['x_train'], f['y_train']  
    x_test, y_test = f['x_test'], f['y_test']
```



- You “pick” the function “shape” – in scikit-learn these are called *classifiers*
- You “fit” the data to the classifier
- You use the classifier to predict

```
from sklearn import tree  
clf = tree.DecisionTreeClassifier()
```

```
clf = clf.fit(x_train, y_train)
```

```
predictions = clf.predict(x_test)
```

# scikit-learn – classifiers

## 1. Supervised learning

- ▶ 1.1. Generalized Linear Models
- ▶ 1.2. Linear and Quadratic Discriminant Analysis
- 1.3. Kernel ridge regression
- ▶ 1.4. Support Vector Machines
- ▶ 1.5. Stochastic Gradient Descent
- ▶ 1.6. Nearest Neighbors
- ▶ 1.7. Gaussian Processes
- 1.8. Cross decomposition

source: [https://scikit-learn.org/stable/user\\_guide.html](https://scikit-learn.org/stable/user_guide.html)

$h(x)$

- ▶ 1.9. Naive Bayes
- ▶ 1.10. Decision Trees
- ▶ 1.11. Ensemble methods
- ▶ 1.12. Multiclass and multilabel algorithms
- ▶ 1.13. Feature selection
- ▶ 1.14. Semi-Supervised
- 1.15. Isotonic regression
- 1.16. Probability calibration
- ▶ 1.17. Neural network models (supervised)




# TensorFlow – concepts


- You “craft” the function “shape” – in TensorFlow this is a *model*

[illegible]

- You “optimize” the model using a *cost/loss function* and *optimizer*



- You use the model to predict



# TensorFlow – concepts

- model – the function shape we construct

$$h(x)$$

- cost/loss function – a function that tells us how bad we are at predicting

$$l(h(x), y) \neq \emptyset$$


- optimizer – method for reducing how bad we are at predicting

# TensorFlow – concepts


- You “craft” the function “shape” – in TensorFlow this is a *model*

[illegible]

- You “optimize” the model using a *cost/loss function* and *optimizer*



- You use the model to predict





# TensorFlow – concepts

- You “craft” the function “shape” – in TensorFlow this is a *model*
- You “optimize” the model using a *cost/loss function* and *optimizer*
- You use the model to predict

```
import tensorflow as tf
from tensorflow import keras
tf.__version__
```

```
# function shape
model = keras.Sequential([
    keras.layers.Dense(128, activation='relu'),
    keras.layers.Dense(10, activation='softmax')
])
```

```
# how to optimize the function
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

```
model.fit(x_train, y_train, epochs=10)
```

```
pred = model.predict(x)
```

# TensorFlow – optimizers

- class **Adadelta**: Optimizer that implements the Adadelta algorithm.
- class **Adagrad**: Optimizer that implements the Adagrad algorithm.
- class **Adam**: Optimizer that implements the Adam algorithm.
- class **Adamax**: Optimizer that implements the Adamax algorithm.
- class **Ftrl**: Optimizer that implements the FTRL algorithm.
- class **Nadam**: Optimizer that implements the NAdam algorithm.
- class **Optimizer**: Updated base class for optimizers.
- class **RMSprop**: Optimizer that implements the RMSprop algorithm.
- class **SGD**: Stochastic gradient descent and momentum optimizer.

# TensorFlow – loss functions

- class **BinaryCrossentropy**: Computes the cross-entropy loss between true labels and predicted labels.
- class **CategoricalCrossentropy**: Computes the crossentropy loss between the labels and predictions.
- class **CategoricalHinge**: Computes the categorical hinge loss between `y_true` and `y_pred`.
- class **CosineSimilarity**: Computes the cosine similarity between `y_true` and `y_pred`.
- class **Hinge**: Computes the hinge loss between `y_true` and `y_pred`.
- class **Huber**: Computes the Huber loss between `y_true` and `y_pred`.
- class **KLDivergence**: Computes Kullback Leibler divergence loss between `y_true` and `y_pred`.
- class **LogCosh**: Computes the logarithm of the hyperbolic cosine of the prediction error.
- class **MeanAbsoluteError**: Computes the mean of absolute difference between labels and predictions.
- class **MeanAbsolutePercentageError**: Computes the mean absolute percentage error between `y_true` and `y_pred`.
- class **MeanSquaredError**: Computes the mean of squares of errors between labels and predictions.
- class **MeanSquaredLogarithmicError**: Computes the mean squared logarithmic error between `y_true` and `y_pred`.
- class **Poisson**: Computes the Poisson loss between `y_true` and `y_pred`.
- class **SparseCategoricalCrossentropy**: Computes the crossentropy loss between the labels and predictions.
- class **SquaredHinge**: Computes the squared hinge loss between `y_true` and `y_pred`.

# TensorFlow – models

```
# function shape
```

```
model = keras.Sequential([  
    keras.layers.Dense(layer_width, activation='relu'),  
    keras.layers.Dense(10, activation='softmax')  
])
```

```
# function shape
```

```
model = keras.Sequential([  
    keras.layers.Reshape((28, 28, 1)),  
    keras.layers.Conv2D(32, (3, 3), activation='relu'),  
    keras.layers.Conv2D(64, (3, 3), activation='relu'),  
    keras.layers.MaxPooling2D(pool_size=(2, 2)),  
    keras.layers.Dropout(0.25),  
    keras.layers.Flatten(),  
    keras.layers.Dense(layer_width, activation='relu'),  
    keras.layers.Dropout(0.5),  
    keras.layers.Dense(10, activation='softmax')  
])
```

```
base_model = tf.keras.applications.MobileNetV2(input_shape=img_shape,  
                                                include_top=False,  
                                                weights='imagenet')
```

```
base_model.trainable = True
```

```
model = tf.keras.Sequential([  
    base_model,  
    tf.keras.layers.GlobalAveragePooling2D(),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```





science

question



research



hypothesis



test



analyze



report

science

question



research



hypothesis


test



analyze



report

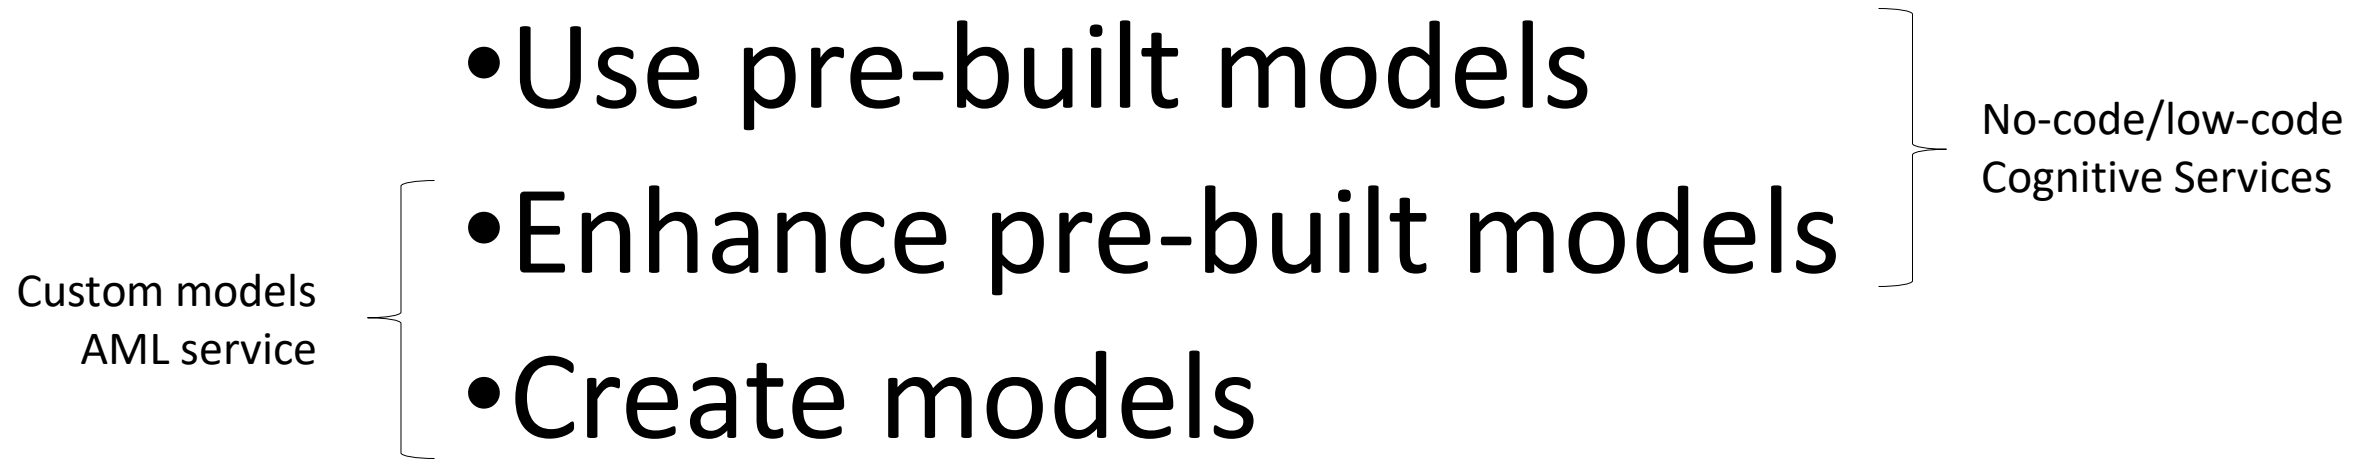




## Three options

- Use pre-built models
- Enhance pre-built models
- Create models

# Three options



# Cognitive Services

... a demo

# Custom Vision

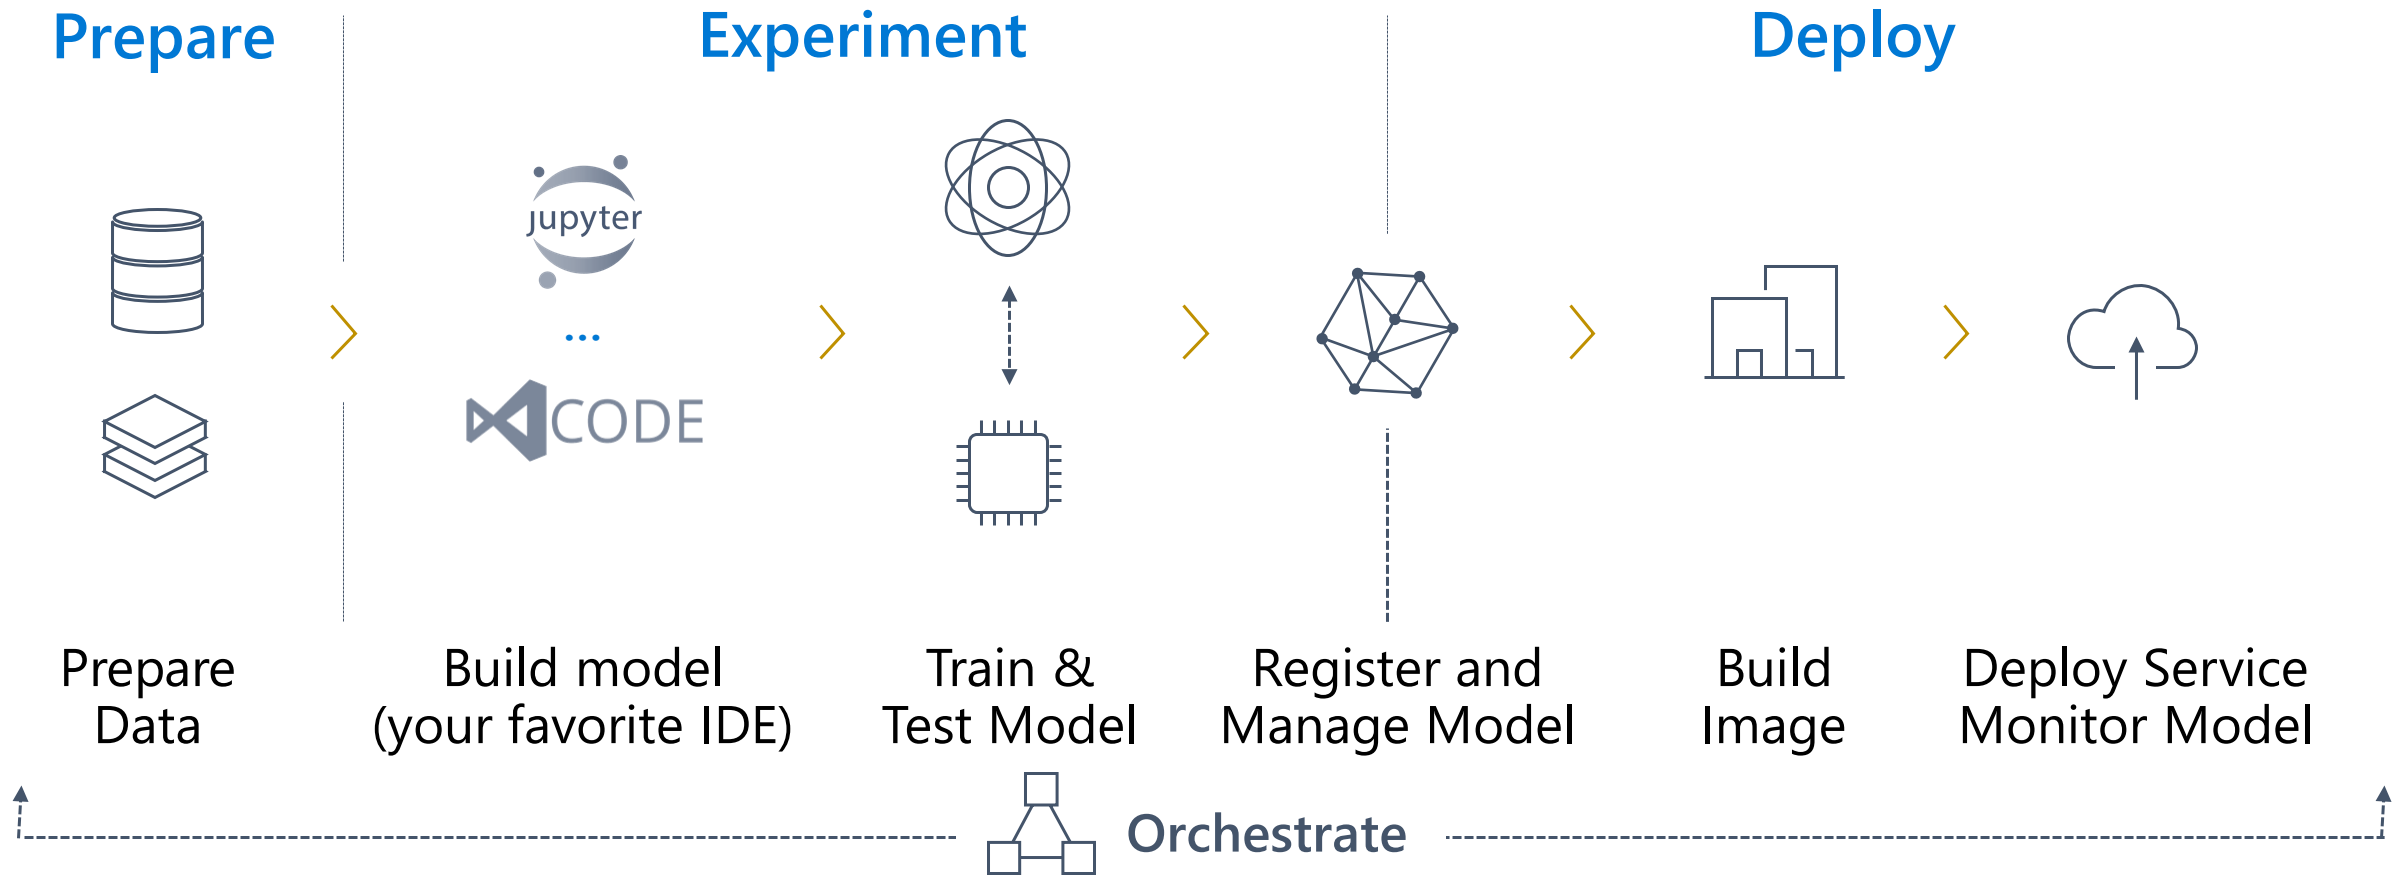
... a demo

# Azure Machine Learning service

... a demo

# Azure Machine Learning

Typical E2E Process






# Azure Machine Learning

- **Datasets** – registered, known data sets
- **Datastores** – Connections to data
- **Compute** – Managed compute
- **Environments** – defined training and inference environments
- **Experiments** – Training runs
- **Pipelines** – Training workflows
- **Models** – Registered, versioned models
- **Endpoints:**
  - Real-time Endpoints – Deployed model endpoints
  - Pipeline Endpoints – Training workflow endpoints



# An Important ML Example



burrito

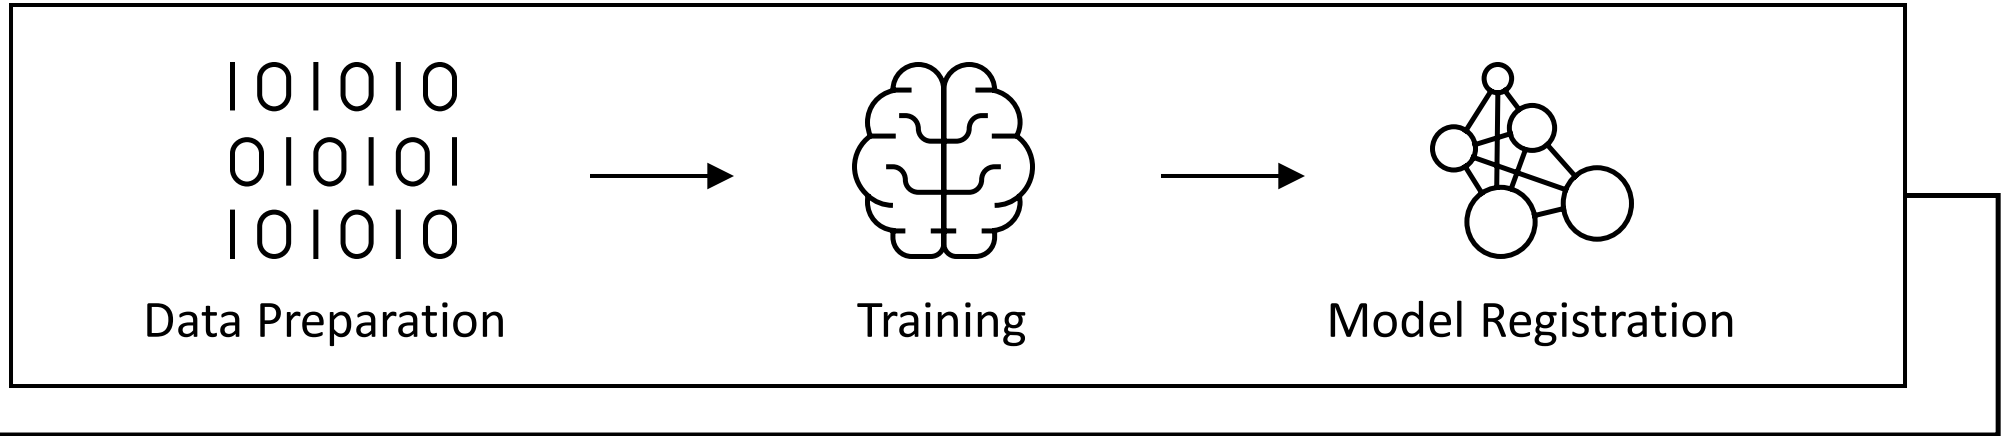


tacos

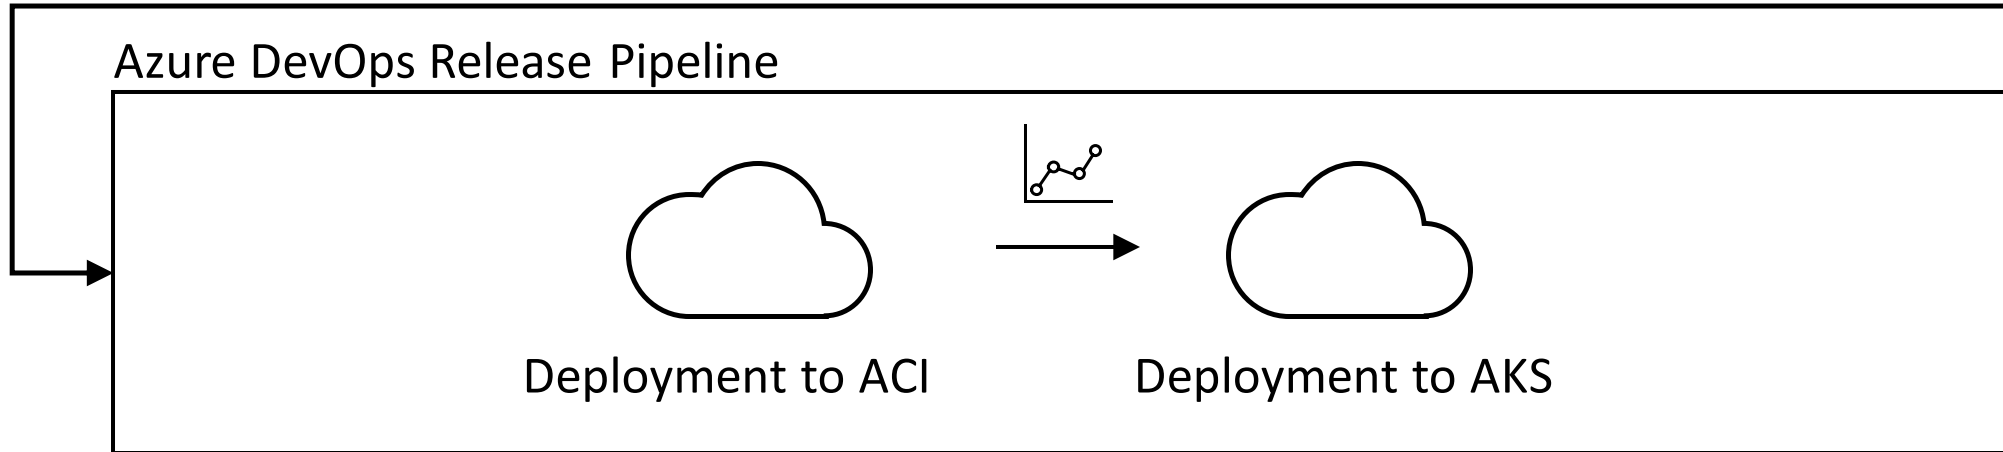


# Complete Pipeline

## AML Pipeline



## Azure DevOps Release Pipeline



# Ethics

- What is my model doing?
- How do I know it is behaving ethically?

# Model Explainability

# Review

- Introduction to Machine Learning
  - a function created by data
- Three-pronged strategy
  - prebuilt models
  - customized prebuilt models
  - custom model
- Cloud based machine learning
  - Azure Machine Learning service
- AI as a Team Sport – managing the AI Dev Lifecycle
  - Azure DevOps <-> Azure Machine Learning service
- Ethics
  - model explainability

Questions?

Seth Juarez  
Microsoft  
@sethjuarez